ORIGAMI
Lab

Containers

QLS 612 - May 2023

Sebastian Urchs - @s_urchs

Fondation Brain Canada Foundation

NIH

NIBIB

HEALTHY BRAINS HEALTHY LIVES

neuro
Montreal Neurological Institute-Hospital

CONP PCNO

Containers?

What we will talk about

1. **Why** are containers useful for researchers
2. **Virtual Machines** (briefly)
3. Using and building containers with **Docker**
4. Using containers on supercomputers with **Singularity**

Why isolate environments

Document your software environment

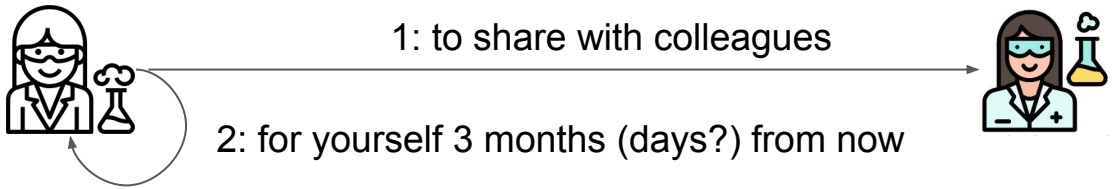

Document your software environment

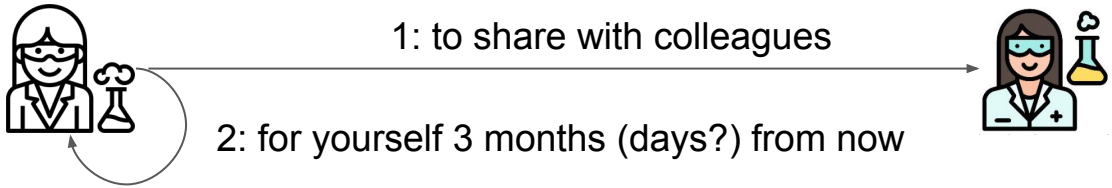

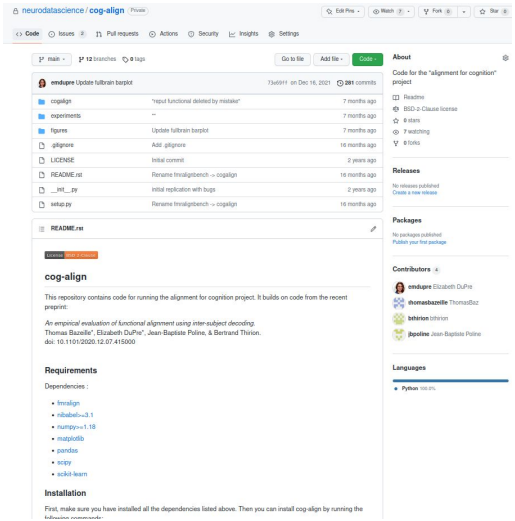

Requirements

Dependencies :

- fmralign
- nibabel>=3.1
- numpy>=1.18
- matplotlib
- pandas
- scipy
- scikit-learn

Installation

First, make sure you have installed all the dependencies listed above.
Then you can install cog-align by running the following commands:

```
git clone https://github.com/neurodatascience/cog-align
cd cog-align
pip install -e .
```

https://github.com/neurodatascience/cog-align

Don’t use the same environment for all projects

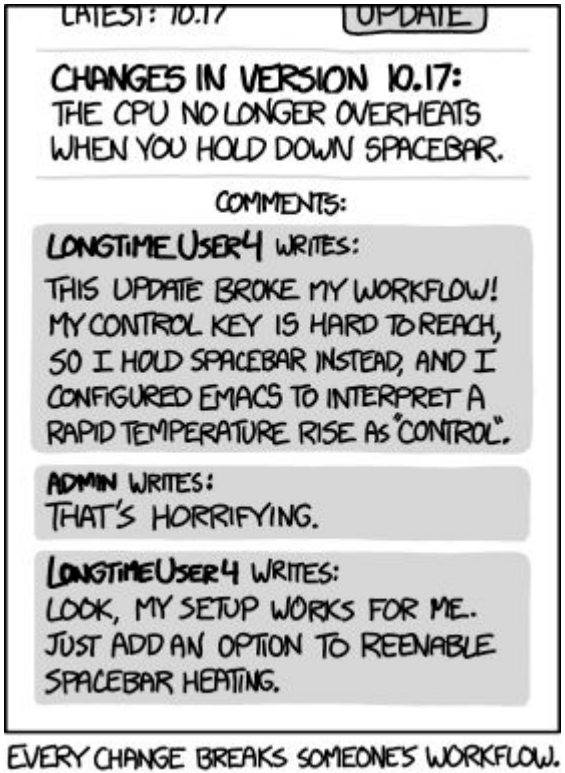

https://xkcd.com/1172/

Changing dependencies may do unexpected things

A. Your updated the dependencies of an existing project

Don’t use the same environment for all projects

Changing dependencies may do unexpected things

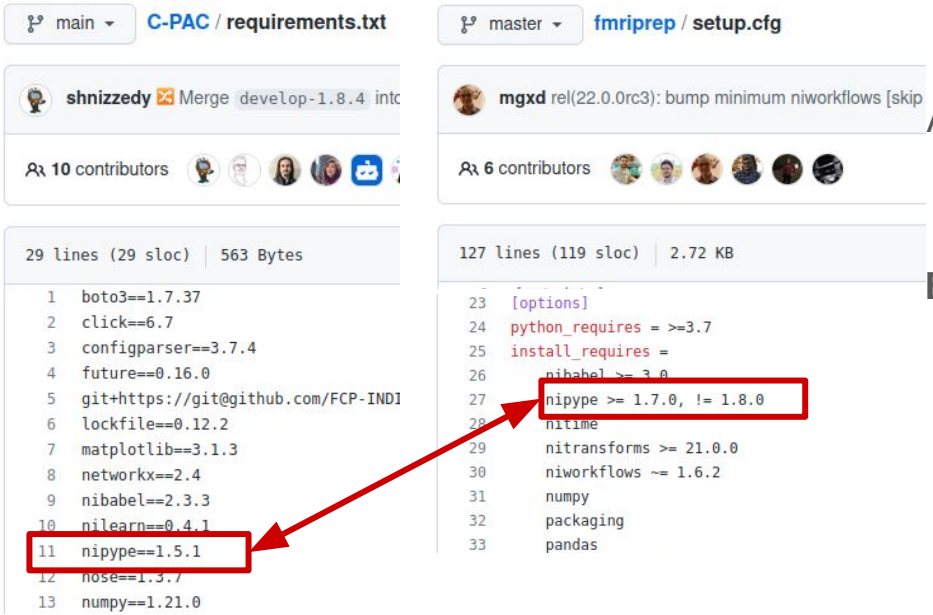

A. Your updated the dependencies of an existing project

B. **Two projects use the same environment but need different versions of some dependencies**

Do not install things into your system Python!

Common installation issues

Installing into the system Python on Linux

On Linux systems, a Python installation will typically be included as part of the distribution. Installing into this Python installation requires root access to the system, and may interfere with the operation of the system package manager and other components of the system if a component is unexpectedly upgraded using `pip`.

On such systems, it is often better to use a virtual environment or a per-user installation when installing packages with `pip`.

```
[surchs@marvin ~]$ which python
/usr/bin/python
[surchs@marvin ~]$ which pip
/usr/bin/pip
[surchs@marvin ~]$
```

https://docs.python.org/3/installing/index.html

Consider: a cake

Home · Cakes · Perfect Cream Cheese Pound Cake

Perfect Cream Cheese Pound Cake

Published by **Sally** on February 18, 2019 - **700 comments**

not sponsored, but I absolutely adore Nordic Ware Bundt pans.
10-12 cups of batter. **This one** is also gorgeous! 🙂

5 **Bake:** Bake the cream cheese pound cake at 325°F (163°C). Hal
the cake with aluminum foil to prevent over-browning.

6 **Cool, then invert:** Let the pound cool for about 2 hours in the
plate and cool completely before serving.

https://sallysbakingaddiction.com/cream-cheese-pound-cake/

Consider: a cake

Home · Cakes · Perfect Cream Cheese Pound Cake

Perfect Cream Cheese Pound Cake

Published by **Sally** on February 18, 2019 - **700 comments**

not sponsored, but I absolutely adore Nordic Ware Bundt pans.
10-12 cups of batter. **This one** is also gorgeous! 🙂

5 **Bake:** Bake the cream cheese pound cake at 325°F (163°C). Hal
the cake with aluminum foil to prevent over-browning.

6 **Cool, then invert:** Let the pound cool for about 2 hours in the
plate and cool completely before serving.

https://sallysbakingaddiction.com/cream-cheese-pound-cake/

Isolate environments to handle different requirements

Why not just python virtual environments?

External Dependencies

fMRIPrep is written using Python 3.8 (or above), and is based on nipype.

fMRIPrep requires some other neuroimaging software tools that are not handled by the Python's packaging system (Pypi) used to deploy the `fmriprep` package:

- FSL (version 6.0.5.1)
- ANTs (version 2.3.3 - NeuroDocker build)
- AFNI (version 22.3.06)
- C3D (version 1.3.0)
- FreeSurfer (version 7.3.2)
- ICA-AROMA (version 0.4.5)
- bids-validator (version 1.8.0)
- connectome-workbench (version 1.5.0)

https://fmriprep.org/en/stable/installation.html

- Not all binary dependencies are on Anaconda
- Not everything is written in Python or R
- Your Operating System (**OS**) also has packages and a package manager
- The same version problems apply to these

Need to capture the (entire) compute environment

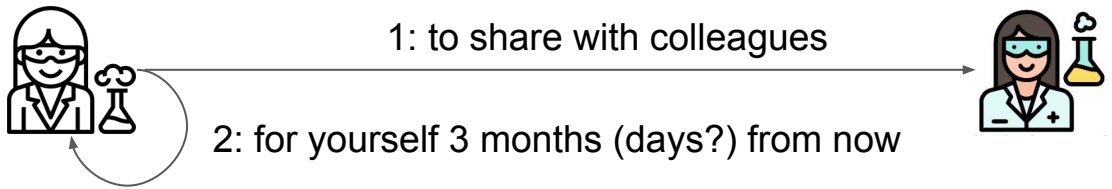

Virtual Machines

How can we isolate different OS environments?

Buy a lot of computers?

Virtual Machines: Let’s simulate the computers

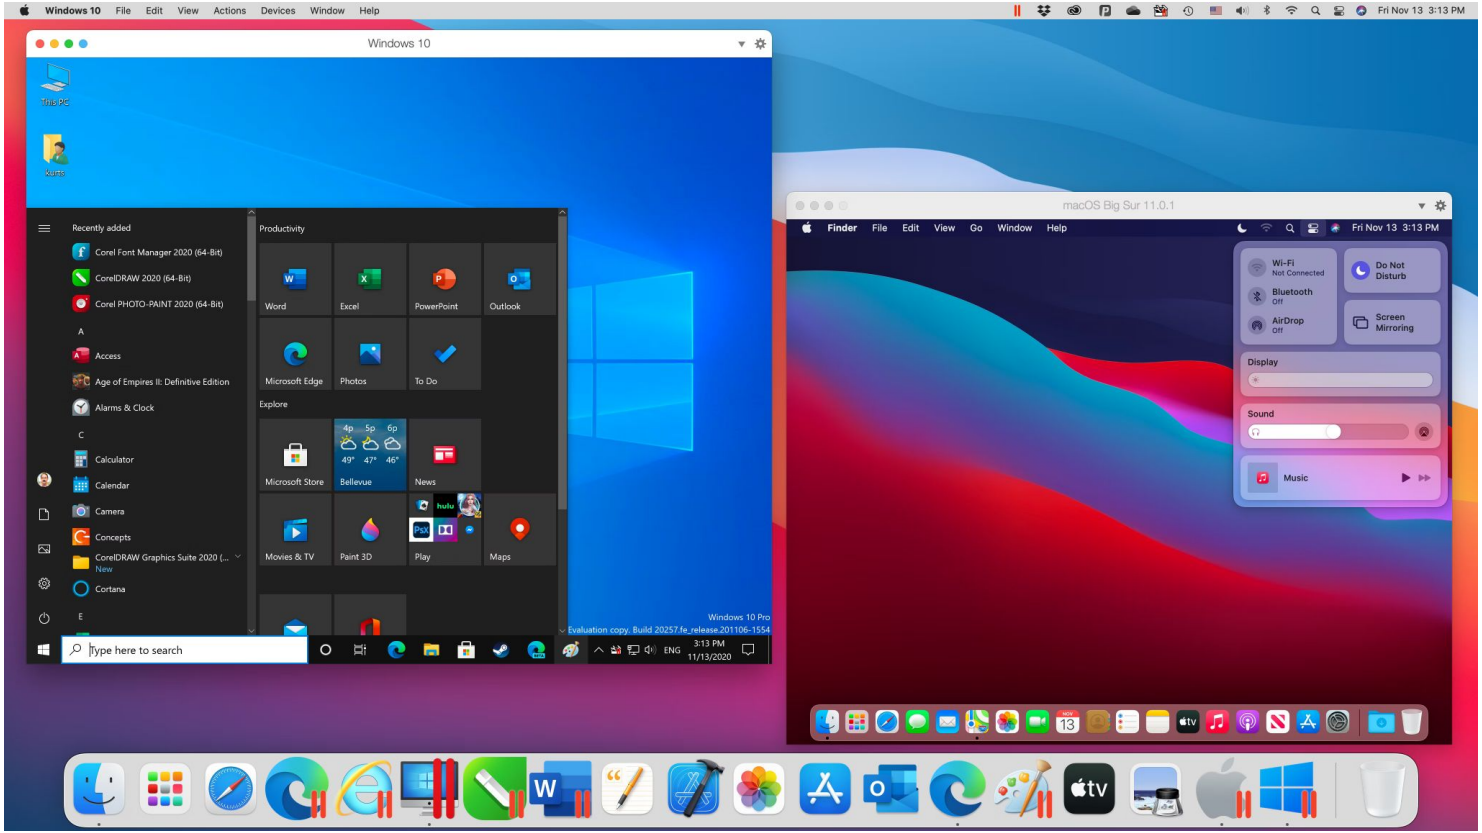

Virtual Machine (VM)

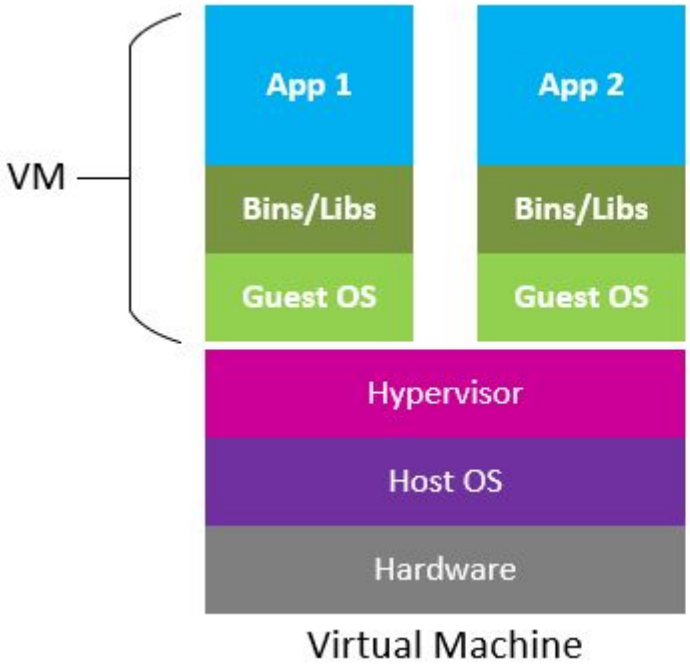

What is it?

- entire computer hardware simulated (**hardware virtualization**)
- you can choose what hardware is exposed (CPU, RAM, hard-drive, …)
- the VM is fully isolated from the Host OS: you get to install any OS inside the VM

How do I make one?

Get Hypervisor / Virtualization software

- VirtualBox (partly OSS)
- VMWare (commercial)
- Hyper-V (Windows)
- …

https://cloudblogs.microsoft.com/opensource/2019/07/15/how-to-get-started-containers-docker-kubernetes/

https://www.virtualbox.org/

Virtual Machine (VM)

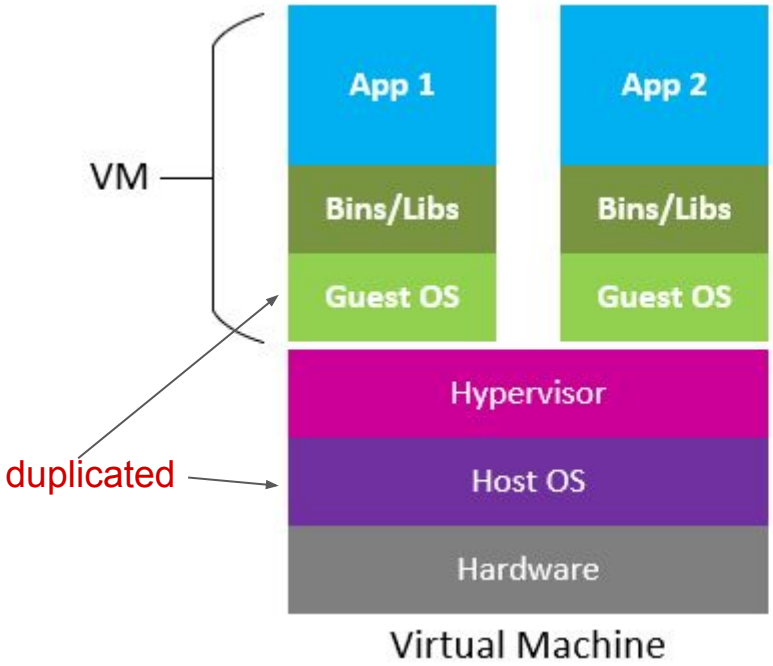

Good:

- can run anything a computer runs (Windows, Linux, Hackintosh)
- can make a snapshot to share with others (Neurodebian used to have a VM)
- good way to test things across many systems
- full system isolation

But:

- doesn't share resources with host -> BIG
- slow to start up, stop, resume
- cumbersome to configure for each project
- duplicates things (every VM needs OS / Kernel)
- no easy way to "get" and use VMs from other people
- you can't use it on a supercomputer

https://cloudblogs.microsoft.com/opensource/2019/07/15/how-to-get-started-containers-docker-kubernetes/

https://www.virtualbox.org/

Containers

Containers: let's all share the same kernel, but in boxes

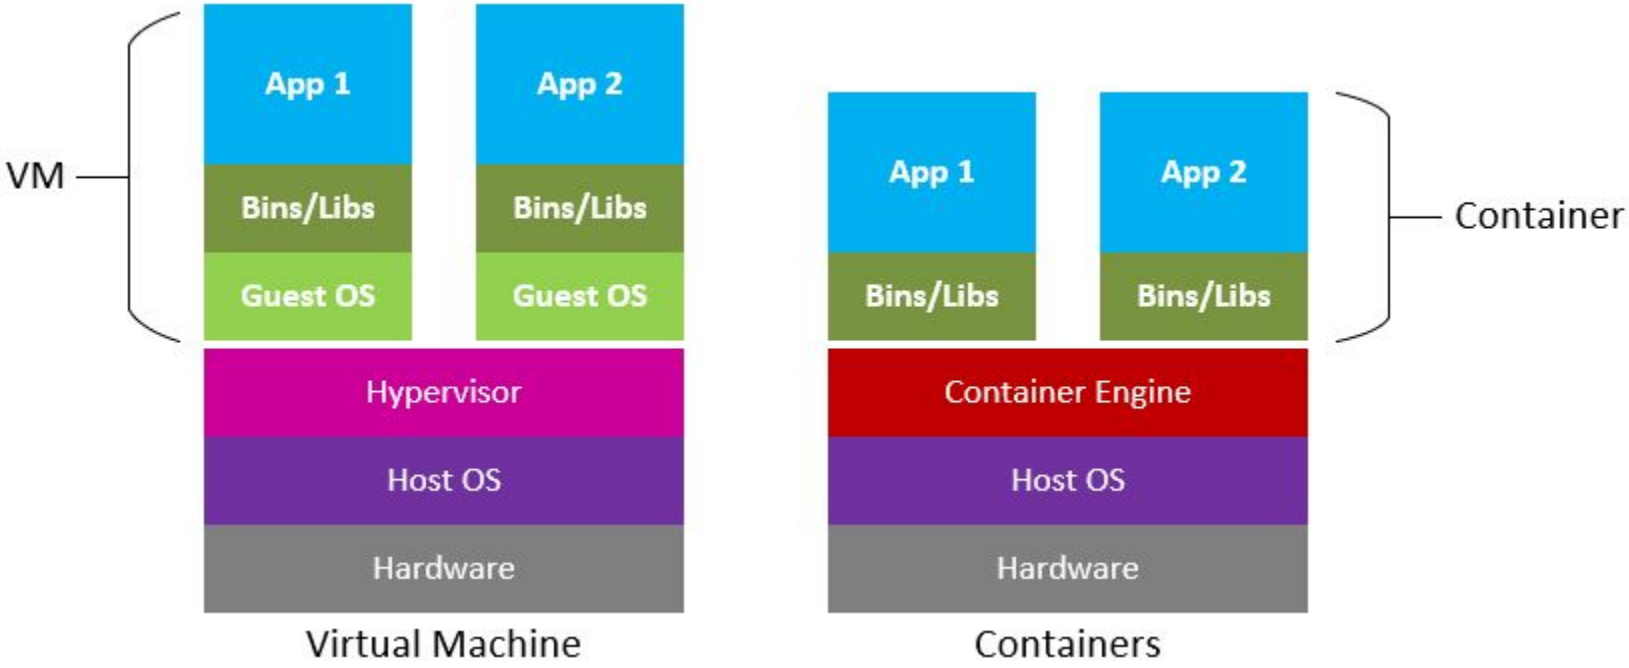

https://cloudblogs.microsoft.com/opensource/2019/07/15/how-to-get-started-containers-docker-kubernetes/

Containers

What are they

- isolated environments sharing the same kernel / OS -> (**OS virtualization**)
- from the inside, a container looks like a separate computer, can't see outside
- within each container, you can have your desired binary and library dependencies

How do I make one

- use a container implementation
- **docker** is the most widely used
- Singularity is used on supercomputers

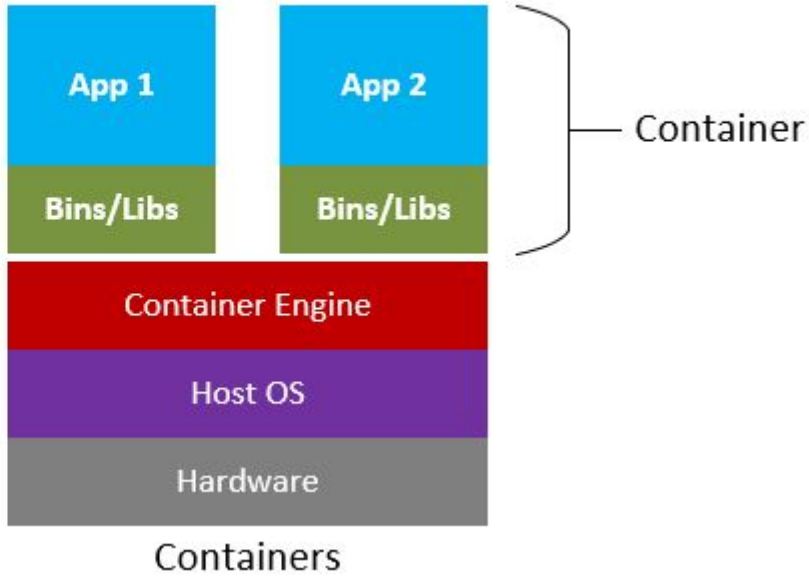

https://cloudblogs.microsoft.com/opensource/2019/07/15/how-to-get-started-containers-docker-kubernetes/

Virtual Machine

Container

What is Docker

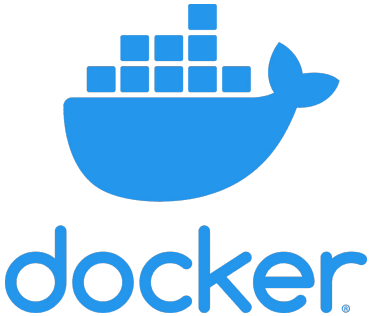

Docker Engine

- a command line program
- gets and builds Docker images and runs Docker containers

Docker Hub

- a website / web service
- a central repository to store and share Docker container images (commercial)
- other container image registries exist

https://docs.docker.com/engine/

https://hub.docker.com/

Docker **image** and **container**: what's the difference

Docker image

- a **read-only** snapshot of an environment
- organized in **layers**
- changing an image adds more layers
- can be stored on Dockerhub or as a file
- images can share identical layers
- can make your own with a **Dockerfile**

Docker container

- a **live instance** of a Docker image
- has a thin writable layer that dies with it
- **one image** can spawn **many containers**

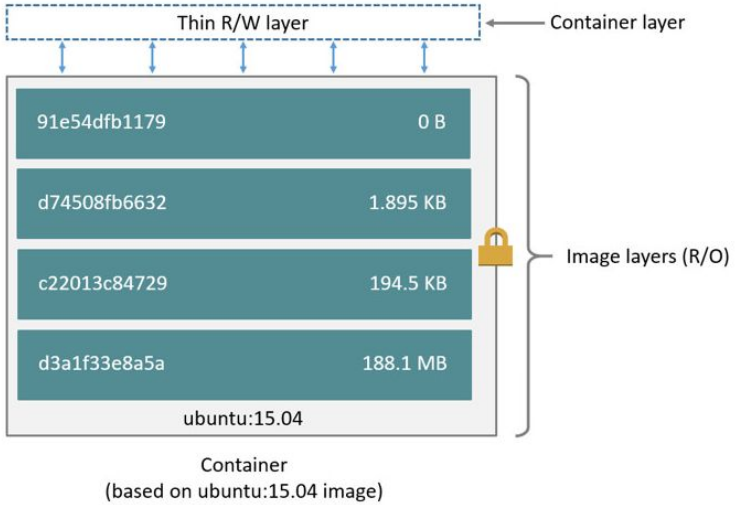

https://docs.docker.com/storage/storagedriver/

How can I get my own Docker container going

Use an **existing image**

- Dockerhub: a repository of docker images https://hub.docker.com/
- You can pull an image with `docker pull`

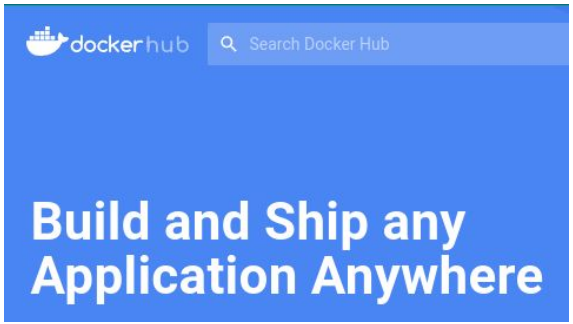

Build your **own image**

- a `Dockerfile` lets you describe the exact image you want to create
- Start from one image and add to it

```
# syntax=docker/dockerfile:1
FROM ubuntu:18.04
LABEL org.opencontainers.image.authors="org@example.com"
COPY . /app
RUN make /app
RUN rm -r $HOME/.cache
CMD python /app/app.py
```

Let’s look at both

Exercise 1: Run a container from Dockerhub

- Find an image we like: https://hub.docker.com/_/hello-world
- Take a look at it
- Pull the image
- Run the image

Exercise 1: Run a container from Dockerhub

- Find an image we like: https://hub.docker.com/_/hello-world
- Take a look at it
- Pull the image
- Run the image

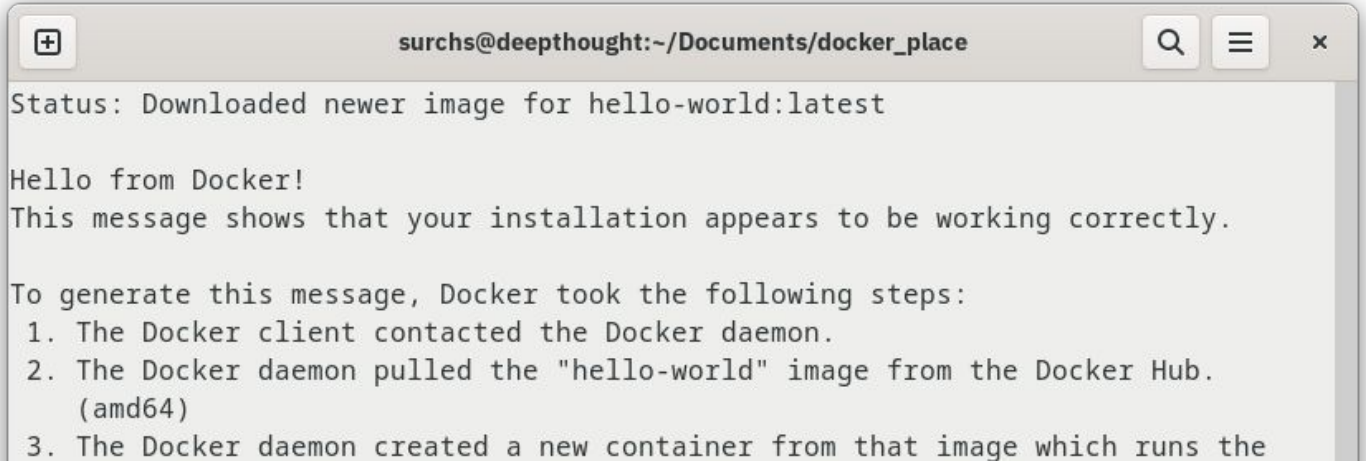

Exercise 1: Run a container from Dockerhub

Summary

- Dockerhub has images
- Image tags are important
- We can
  - retrieve images with `docker pull`
  - create and immediately run a container from an image with `docker run`

Let’s find a more useful image

Exercise 2: work with a container that has conda

I don’t have conda installed on my system. But there is a Docker image. Let’s try!

- Find a Docker image: https://hub.docker.com/r/continuumio/miniconda3/
- Pick a tag, then pull the image `$ docker pull continuumio/miniconda3:22.11.1-alpine`
- Run it

Exercise 2: work with a container that has conda

I don’t have conda installed on my system. But there is a Docker image. Let’s try!

- Find a Docker image: https://hub.docker.com/r/continuumio/miniconda3/
- Pick a tag, then pull the image `$ docker pull continuumio/miniconda3:22.11.1-alpine`
- Run it

Oh, nothing happened?

- Let’s run it interactively to take a look inside

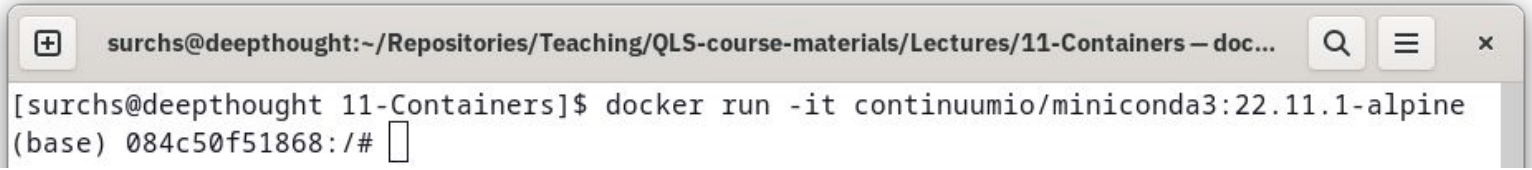

https://github.com/ContinuumIO/docker-images/blob/master/miniconda3/alpine/Dockerfile

Looking around inside a container

- no conda on my machine
- starting container changes the look of my terminal and the name of my computer
- inside of the container I have access to conda

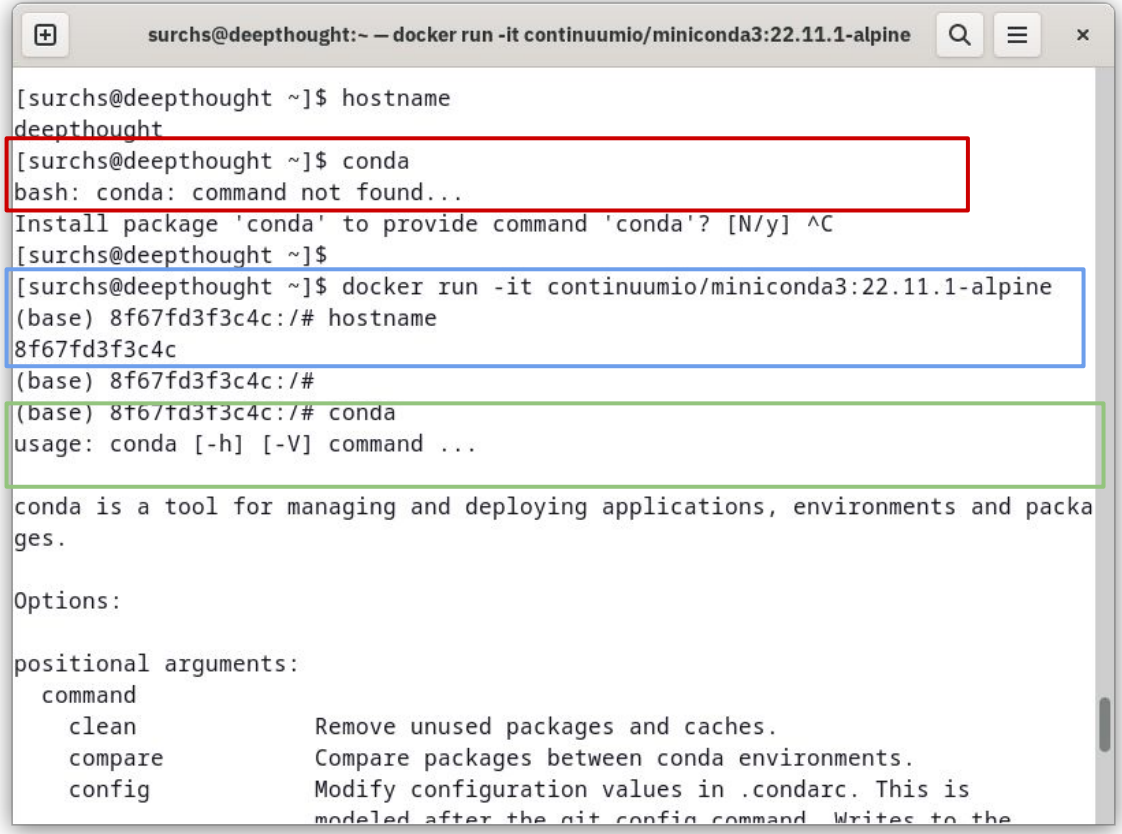

By default the container doesn’t see files on the host …

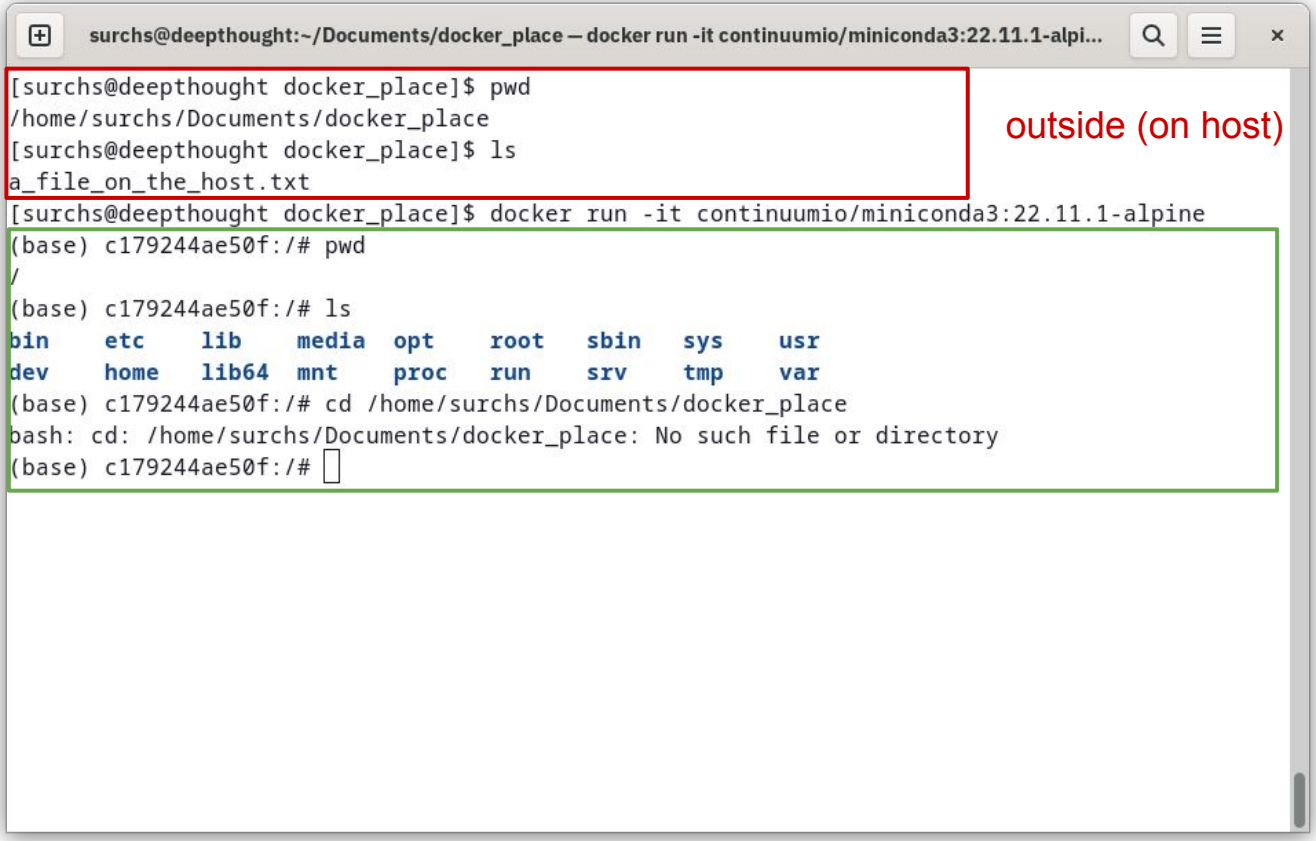

… and the host can’t see files on the container …

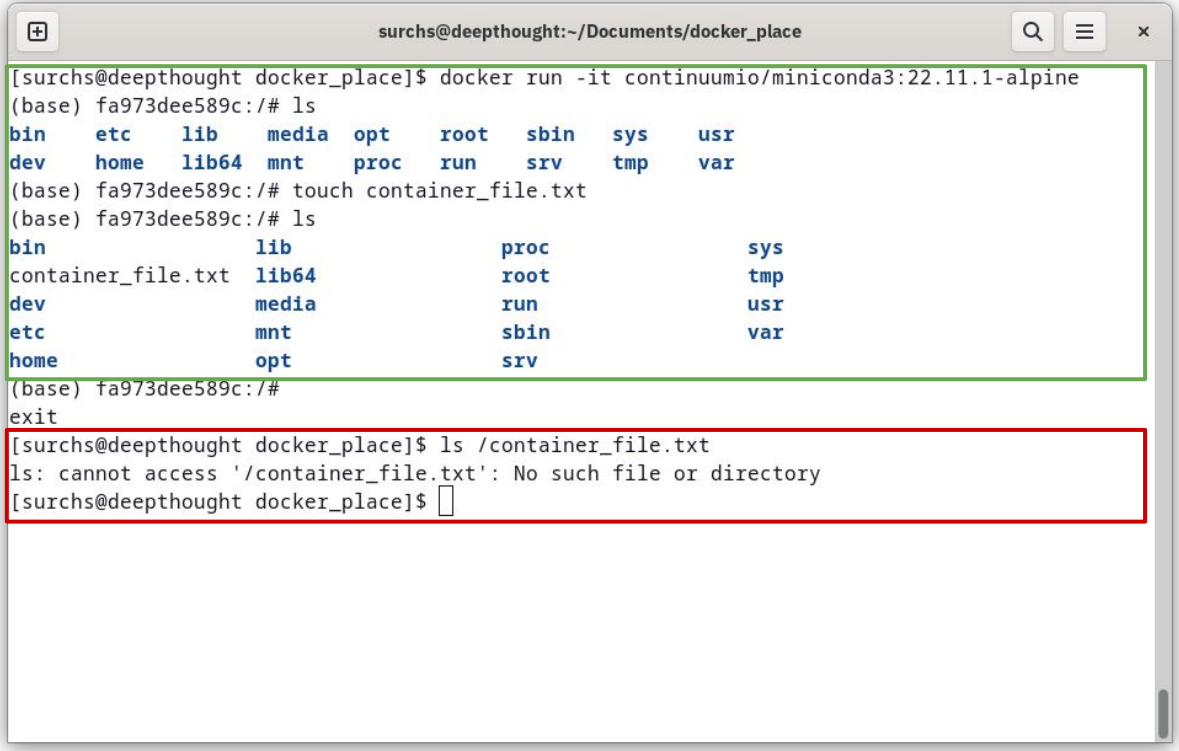

inside
(in container)

outside (on host)

… and files written to a container are tied to it!

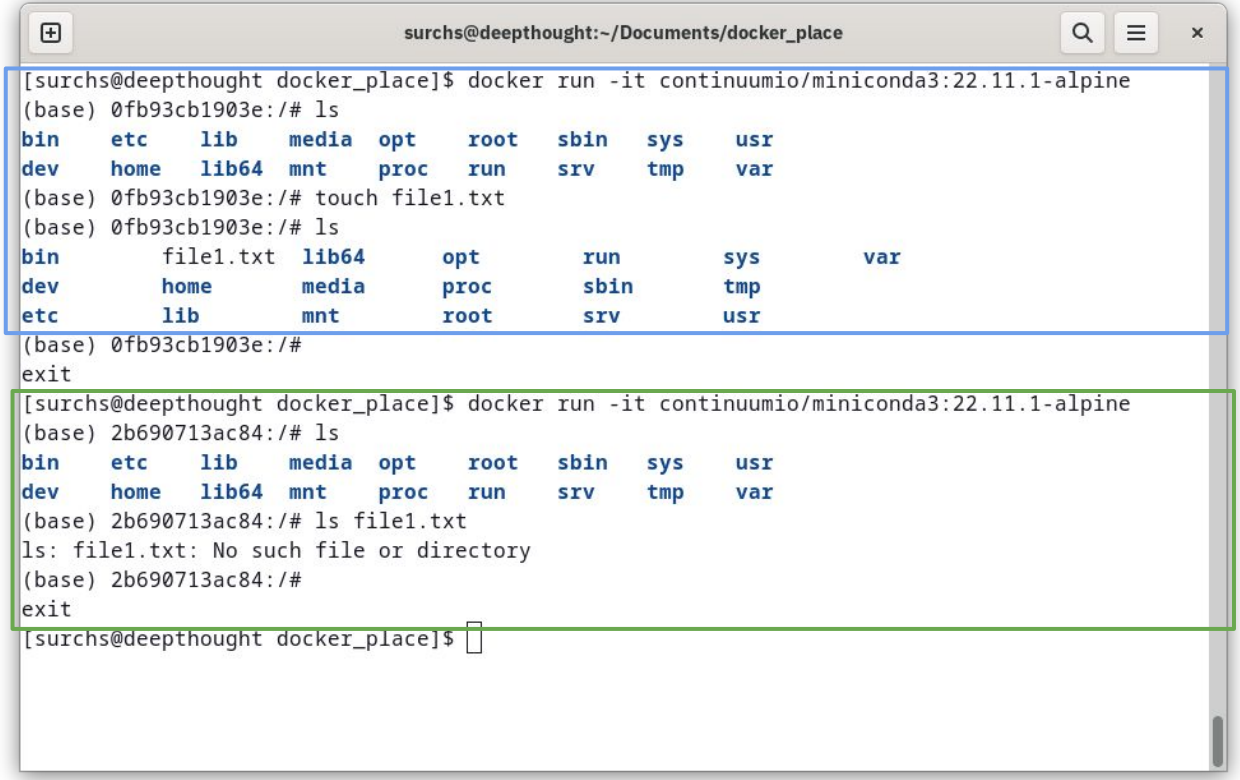

first container instance

second container instance

-> Don’t write anything (worth keeping) to a container

How do we share files between host and container?

- Bind-mount a path on the host to a path in the container

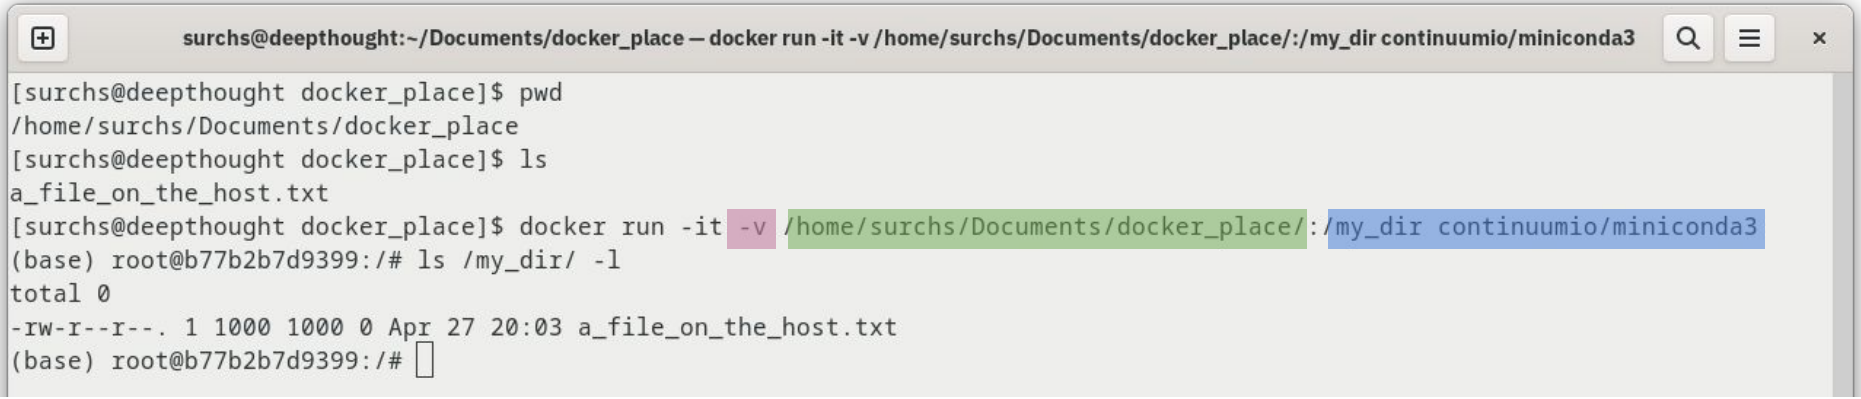

- The bind-mount will be created in the container, even if it exists already!
- Make sure you provide a path (starts with / or ./) - or use `--mount`

https://docs.docker.com/storage/bind-mounts/

Exercise 2: work with a container that has conda

Summary

- We can connect to an interactive shell in a container with `docker run -it`
- By default the container cannot see or write the filesystem of the host
- We can “mount” a path on the host into the container with

```
docker run -v /host/path:/container/path OR
docker run --mount type=bind,source=/host/path,target=/container/path
```

- It’s a bad idea to write into the container directly

So how do I make persistent changes to the image?

Exercise 3: Containerize a classifier

- I want to run the classification script from `08-Machine_Learning_1`
- but I did not setup all of the requirements to run it

surchs@deepthought:~/Repositories/Teaching/QLS-course-materials/Lectures/11-Co...

File: ABIDE_classification.py

```
1  ## Imports
2  from nilearn import datasets
3  from nilearn.connectome import ConnectivityMeasure
4  import pandas as pd
5  import numpy as np
6  import argparse
7  from sklearn import preprocessing
8  from sklearn.model_selection import train_test_split
9  from sklearn.linear_model import LogisticRegression
10 from sklearn.ensemble import RandomForestClassifier
11 from sklearn.metrics import precision_recall_fscore_support
12 from sklearn.metrics import confusion_matrix
13
14
15 def extract_connectome_features(func_data, measure):
16     ''' A function to calculate connnectome based on timeseries data an
   d similarity measure
17     '''
18     connectome_matrix = measure.fit_transform([func_data])[0]
19     tril_idx = np.tril_indices(len(connectome_matrix),k=-1)
```

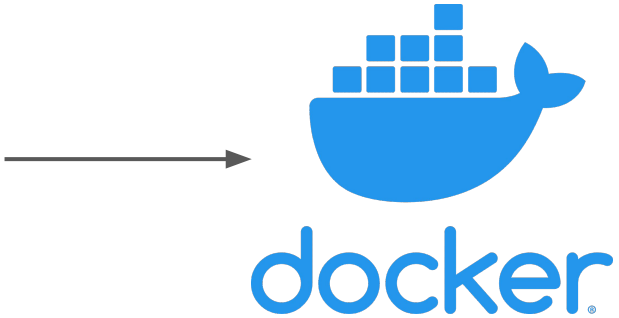

Exercise 3: Containerize a classifier

1. Start **from** an existing image
2. Define **changes** on top of that with a **Dockerfile**
3. Build a new image using the **docker build** command

surchs@deepthought:~/Repositories/Teaching/QLS-course-materials/Lectures/11-Container...

```
FROM python:3.9-slim

RUN pip install --no-cache-dir  numpy nilearn scikit-learn matplotlib seaborn pandas

WORKDIR /predict
COPY ./ABIDE_classification.py /predict

CMD ["/bin/bash"]
```

surchs@deepthought:~/Repositories/Teaching/QLS-course-materials/Lectures/11-Containers

```
[surchs@deepthought 11-Containers]$ docker build -t my_qls_img .
[+] Building 1.8s (10/10) FINISHED
 => [internal] load .dockerignore                                    0.0s
```

Exercise 3: Containerize a classifier

4. We have **built our own docker image and stored it locally** on the computer

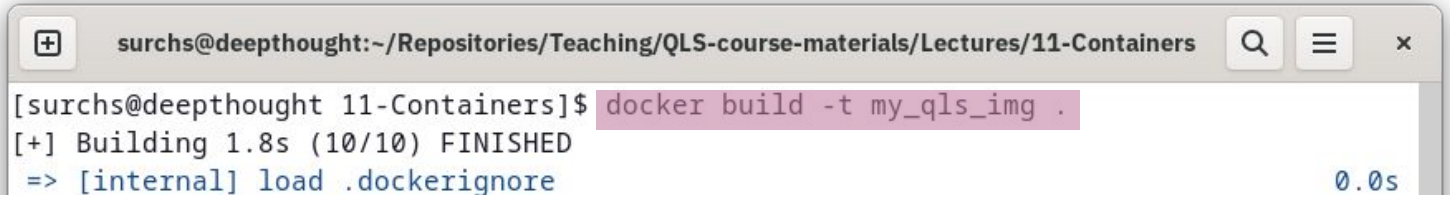

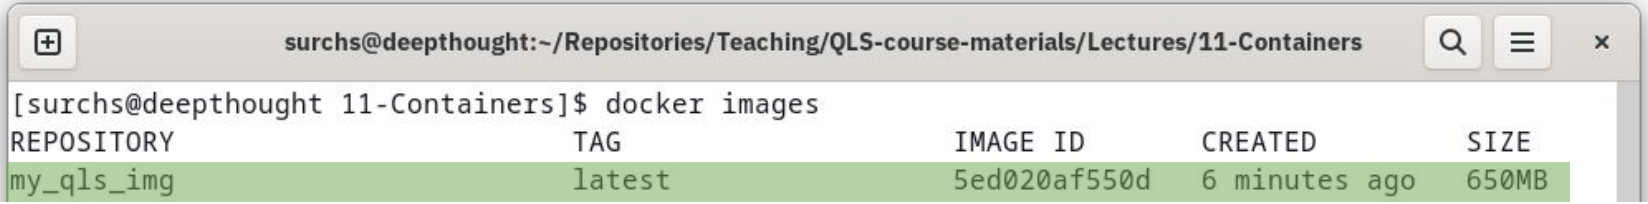

Run the classifier script

1. **Run the container and attach** an interactive shell
2. Inside the container, **run the classification script** like we would on the host

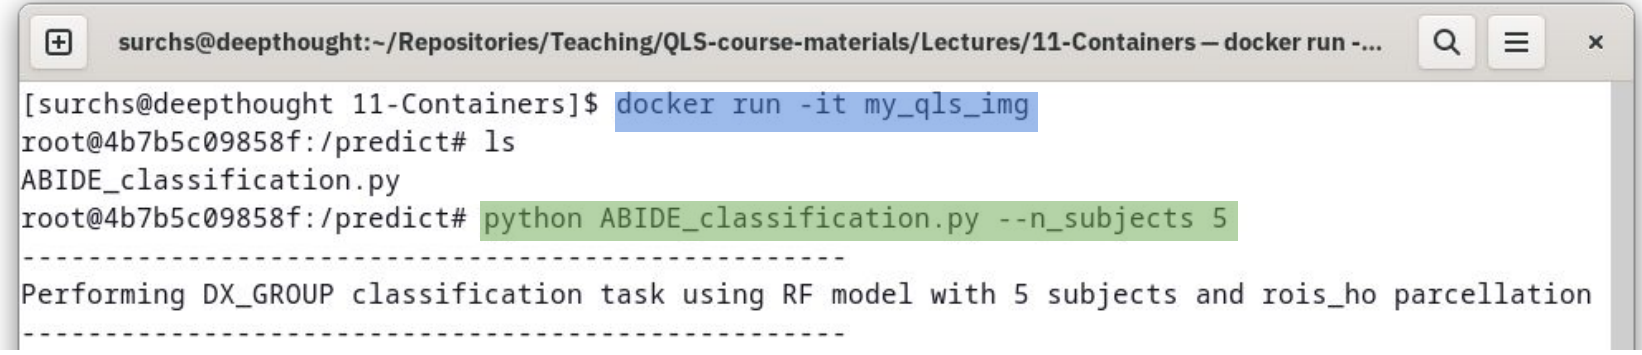

3. We can also **run the container** and then **run the script as an argument**

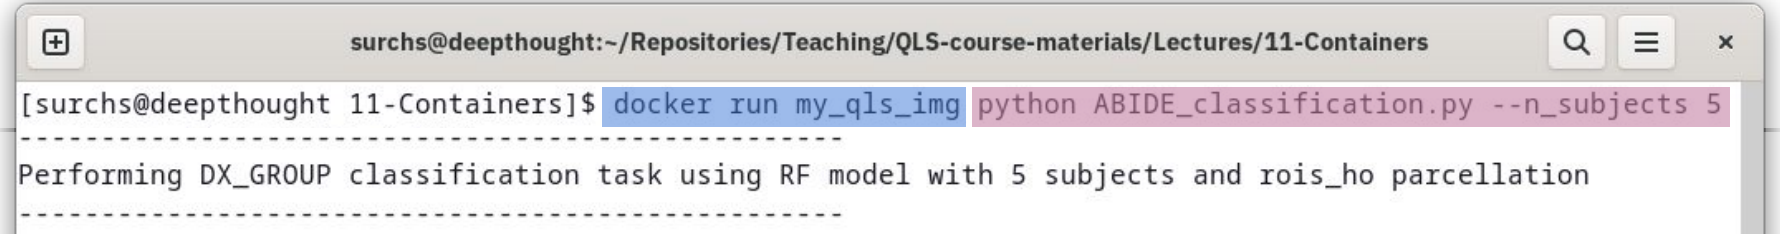

I keep writing the same command, what if I didn't have to?

1. The `CMD` **keyword in the Dockerfile** defines a default command
2. Unless the user provides their own command, the default CMD is run

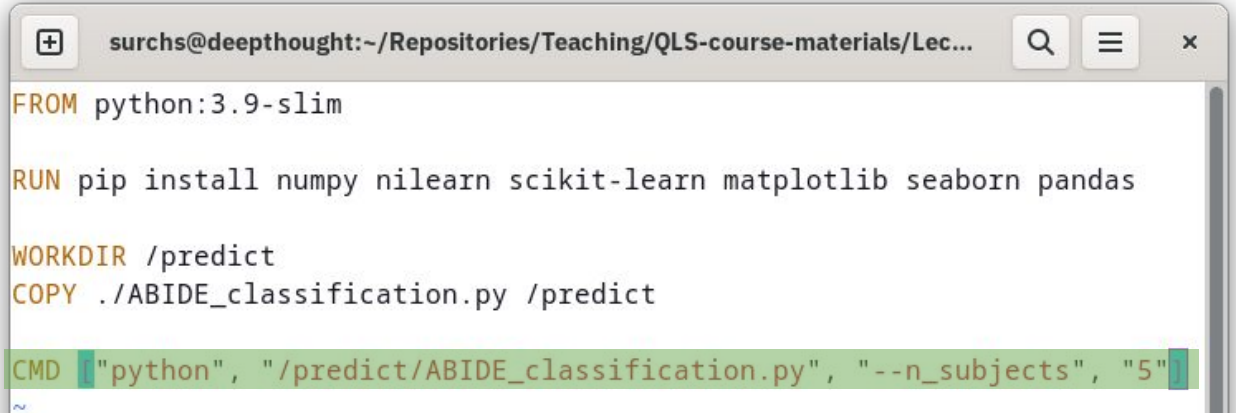

3. To reflect the changes in the Dockerfile, I have to **rebuild** my image

Exercise 3: Containerize a classifier

First summary:

- A `Dockerfile` lets us specify exactly what environment we want in a container
- `CMD` defines a default command that runs if the user doesn't provide one
- We can build our own docker image on top of an existing one
- `docker build` creates the specified docker image on our host computer

Exercise 3: Containerize a classifier

First summary:

- A `Dockerfile` lets us specify exactly what environment we want in a container
- `CMD` defines a default c ser doesn't provide one
- We can build our ow existing one
- `docker build` crea ge on our host computer

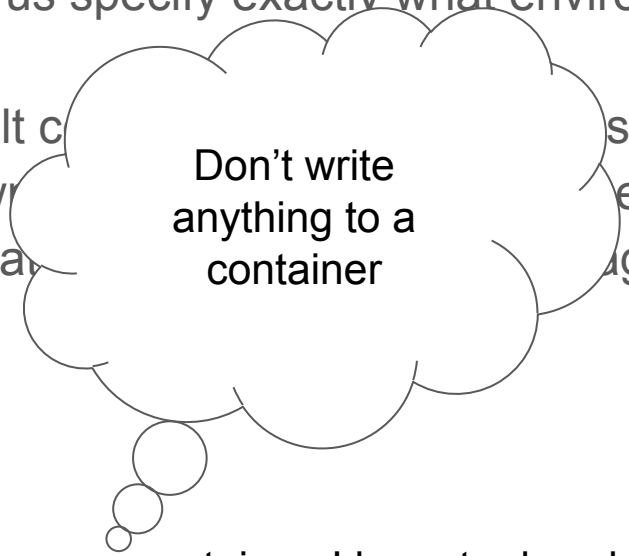

Every time I run a new container, I have to download my data again …

Loading or storing data in a container is not great

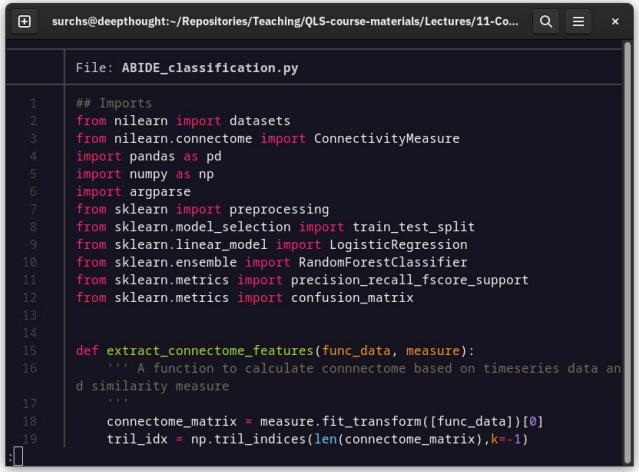

not persistent, not shareable with other container, slow, …

Containers are for processes, not data

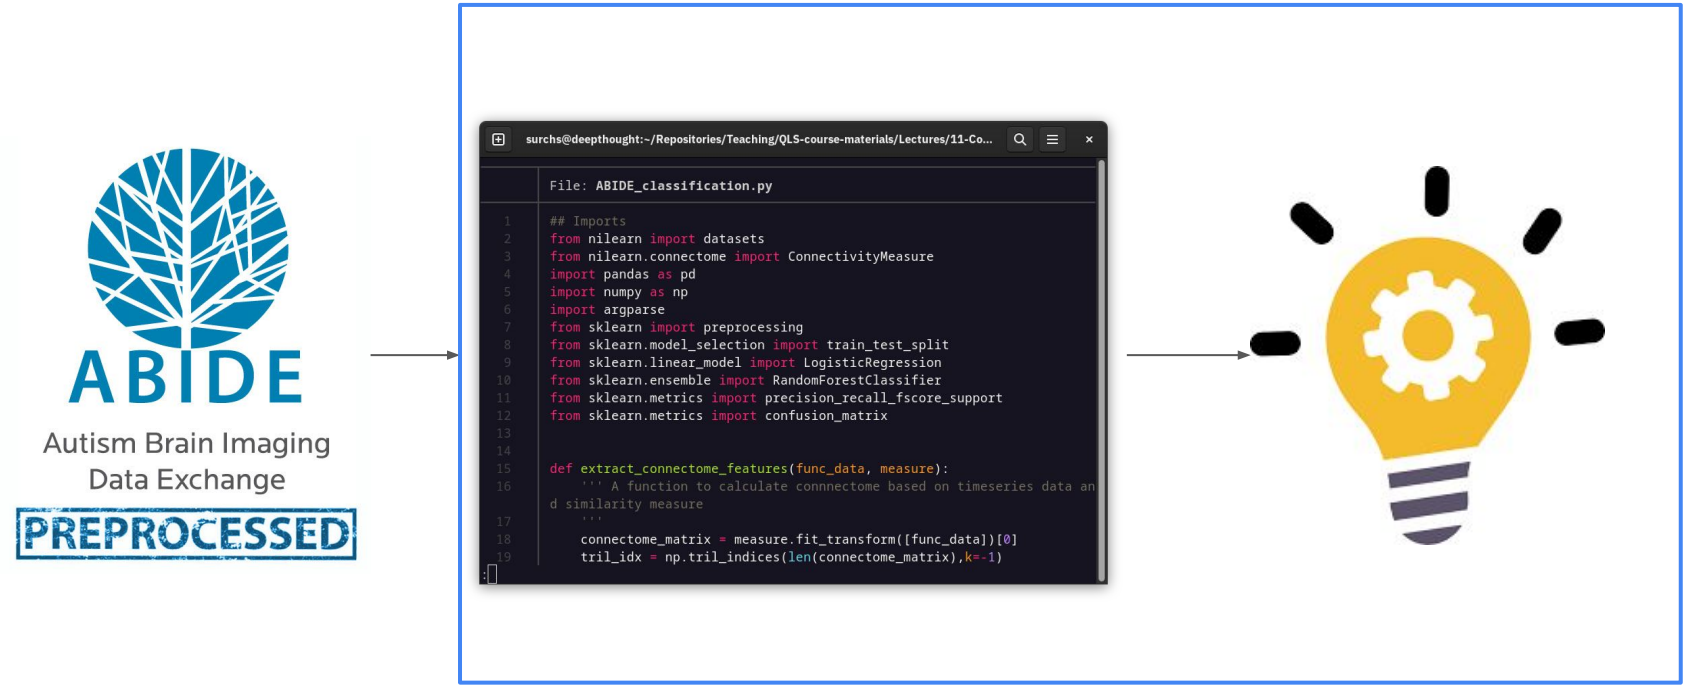

-> let's bind-mount the data into the container from our host machine

Provide the downloaded data via bind-mount

1. Bind the **path to the data on your computer** to a **path in the container**
2. Provide the path in the container **to the classifier script**
3. The script inside the container only sees the mounted path

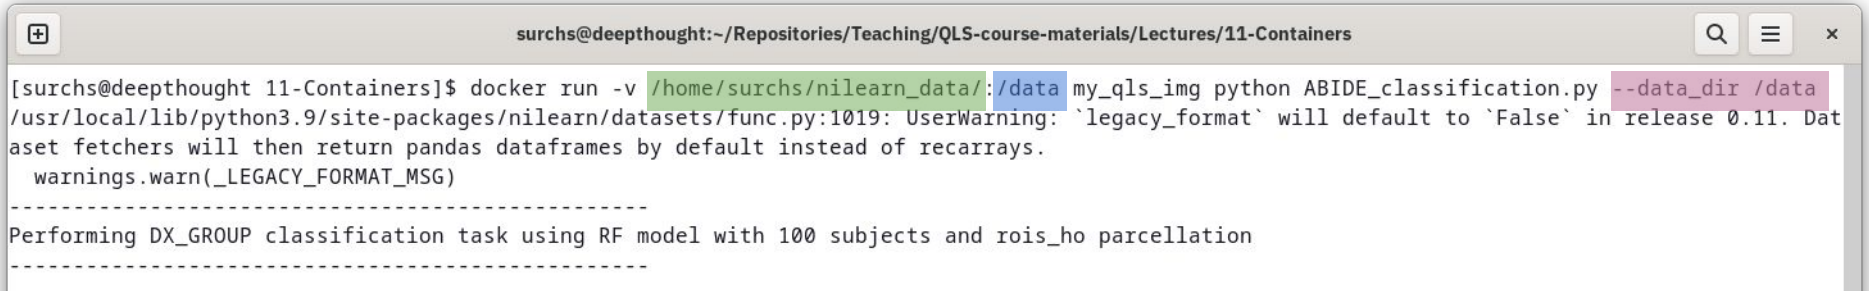

Provide the downloaded data via bind-mount

1. Bind the **path to the data on your computer** to a **path in the container**
2. Provide the path in the container **to the classifier script** as a parameter
3. The script inside the container only sees the mounted path

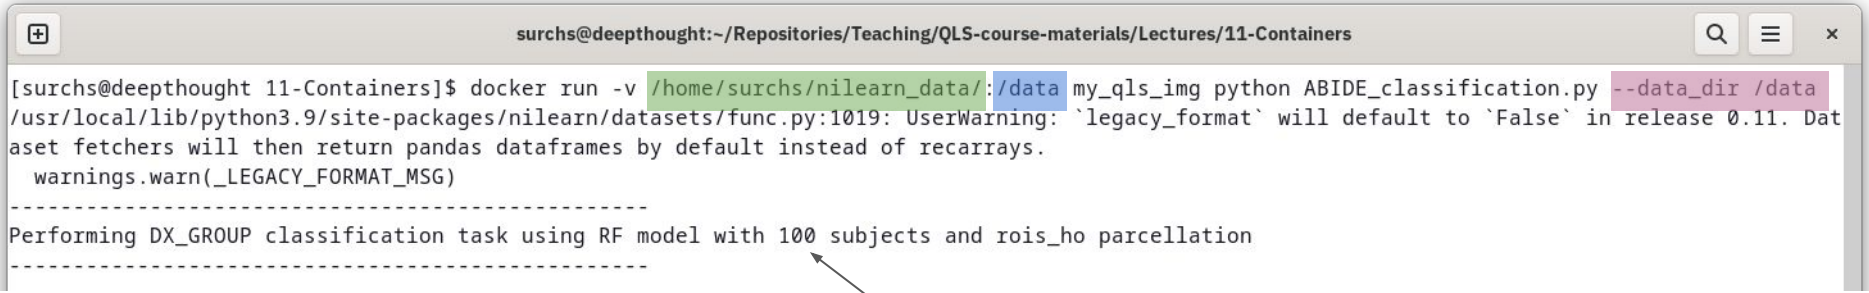

Bonus: why is the classifier running with 100 subjects now?

Exercise 3: Containerize a classifier

- We have a `Dockerfile` that describes our desired environment
- We have built a docker image that contains all dependencies and the code
- We have a default command so a user can just run a container
- **-> We have containerized the classifier**

Exercise 3: Containerize a classifier

- We have a `Dockerfile` that describes our desired environment
- We have built a docker image that contains all dependencies and the code
- We have a default command so a user can just run a container
- **-> We have containerized the classifier**

What if I'd like to make a change to the code?

Exercise 3: Containerize a classifier

- We have a `Dockerfile` that describes our desired environment
- We have built a docker image that contains all dependencies and the code
- We have a default command so a user can just run a container
- **-> We have containerized the classifier**

What if I'd like to make a change to the code?

```
$ docker build .
$ docker build .
$ docker build .
$ docker build .
$ docker build .
$ docker build .
```

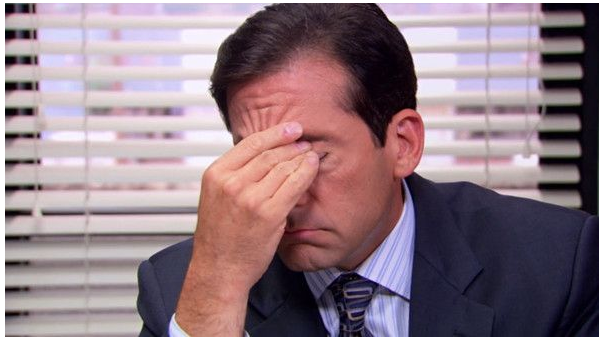

Exercise 3: Containerize a classifier

- We have a `Dockerfile` that describes our desired environme
- We have built a docker image that contains al
- We have a default command so a user can i
- **-> We have containerized the classifie**

The bind-mount will be created in the container, even if it exists already!

What if I'd like to make a change to t

```
$ docker build .
  $ docker build .
    $ docker build .
      $ docker build .
        $ docker build .
          $ docker build .
```

Exercise 3: Containerize a classifier

- We have a `Dockerfile` that describes our desired environme
- We have built a docker image that contains al
- We have a default command so a user can i
- **-> We have containerized the classifie**

The bind-mount will be created in the container, even if it exists already!

What if I'd like to make a change to t

```
$ docker build .
  $ docker build .
    $ docker build .
      $ docker build .
        $ docker build .
          $ docker build .
```

Exercise 3: Keep developing the containerized classifier

- I can bind-mount my local source code over the copied code in the container
- I can keep editing my files locally and test them in the container
- When I build a new image, the most recent files will be copied into the image

🎉

Exercise 3: Containerize a classifier

Summary:

- Dockerfile + docker build create a **controlled image** for our code
- bind-mounts allow us to **expose data** to the container and store results
- We can combine **local editing** of source files and
  the **controlled environment** of the container by bind-mounting our source code
- -> We can use the **same environment** for development, analysis, and sharing

Exercise 3: Containerize a classifier

Summary:

- Dockerfile + docker build create a **controlled image** for our code
- bind-mounts allow us to **expose data** to the container and store results
- We can combine **local editing** of source files and
  the **controlled environment** of the container by bind-mounting our source code
- -> We can use the **same environment** for development, analysis, and sharing

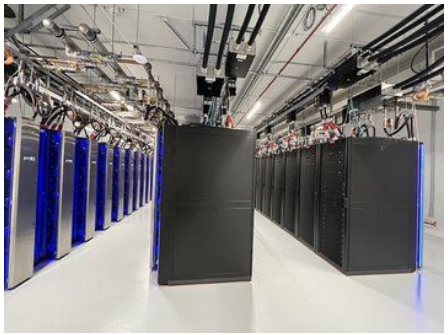

Can I use my new container on a supercomputer?

Singularity

Singularity is the container solution for HPCs

On shared systems (like a supercomputer), you shouldn't / can't use Docker

- Docker isn't as isolated as a VM
- By default you run docker with root privileges and are root inside a container
- A malicious actor can escalate privileges and "break out" of a container

Apptainer (formerly Singularity) **is a container solution in these cases**

Singularity[1] is open source software created by Berkeley Lab:

- as a **secure way** to use Linux containers on Linux multi-user clusters,
- as a way to enable users to have **full control of their environment,** and,
- as a way to **package scientific software** and deploy such to *different* clusters having the *same* architecture.

i.e., it provides **operating-system-level virtualization** commonly called *containers*.

https://apptainer.org/docs/user/main/ and https://apptainer.org/docs/user/main/security.html

https://docs.alliancecan.ca/wiki/Singularity

Build images with Docker, run them with Singularity

1. Pull an image from Dockerhub and create a local **S**ingularity**I**mage**F**ile

```
$ apptainer pull docker://sylabsio/lolcow
```

2. Run the Singularity File using **apptainer run**

```
$ apptainer run lolcow_latest.sif
 _______________________________
< Mon Aug 16 13:01:55 CDT 2021 >
 -------------------------------
        \   ^__^
         \  (oo)\_______
            (__)\       )\/\
                ||----w |
                ||     ||
```

https://apptainer.org/docs/user/main/quick_start.html#

https://docs.alliancecan.ca/wiki/Singularity

Container Summary

- Docker **images** are **read-only snapshots**, you can find them on Dockerhub
- Images have tags (or version), that you should specify when pulling
- A Docker **container** is an isolated, **live instance** of an image
- Docker **containers** have their own file system, but this is **not persistent**
- With **volumes** we can **expose directories** on the host to the container
- We can **build** our own images on top of existing ones using a **Dockerfile**
- Use **Singularity/Apptainer** to run Docker images on a compute cluster

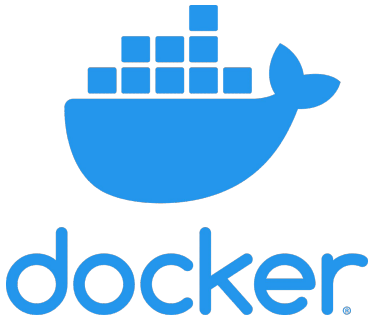

Docker engine

```
# our base image
FROM alpine:3.5

# Install python and pip
RUN apk add --update py2-pip

# upgrade pip
RUN pip install --upgrade pip

# install Python modules needed by the Python app
COPY requirements.txt /usr/src/app/
```

Dockerfile

Docker image registry

There are tools to help make Dockerfiles

Welcome to Neurodocker!

Neurodocker is a command-line program that generates custom Dockerfiles and Singularity recipes for neuroimaging and minifies existing containers. Its purpose is to make it easier for scientists (and others) to easily create reproducible computational environments.

(This requires having Docker installed)

```
neurodocker generate docker --pkg-manager apt \
    --base-image neurodebian:buster \
    --ants version=2.3.4 \
    --miniconda version=latest conda_install="nipype notebook" \
    --user nonroot
```

https://www.repronim.org/neurodocker/user_guide/index.html

Additional Resources

- Docker tutorial https://github.com/docker/labs
- Neurohackweek container course https://neurohackweek.github.io/docker-for-scientists/
- The Turing Way on reproducible research environments https://the-turing-way.netlify.app/reproducible-research/renv.html