

# Introduction to High Performance Computing (HPC)

QLS 612 - 2023  
Brent McPherson, PhD



HBHL  
NEUROHUB

# The goals of this lecture are to:

- Describe the structure of HPCs and how they differ from other traditional computers and common approaches.
  - Laptops / Desktops
- Introduce the concepts of using High Performance Computing (HPCs).
- How to begin effectively using HPC clusters in your work.

*This will reference concepts covered in the “Introduction to the Terminal and Bash” and “Containerization” lectures.*

# Modern NeuroData Science Requires Computers

- No part of our research can progress without access to useful computing resources.
  - Data is acquired and stored on computers.
  - Data preparation and curation is performed on computers.
  - Analysis and modeling require computers.
  - Visualization and reporting of results require computers.
  - The distribution and reporting of results require computers.
- If at any point our work cannot continue in a digital space, it has effectively ended.
- In order to maximize our ability to perform innovative work in a field dependent on modern computational advances, a useful understanding of these machines is necessary.



## Big Computer on Campus: Universities Graduate to AI Super Systems

Colleges across the U.S. and Europe are embracing high-performance systems to accelerate machine learning across all fields of research and education.

August 17, 2021 by [CHERYL MARTIN](#)



# The Concepts of High Performance Computing (HPC)

# Working on a personal computer (PC)

- Everyone is familiar with this.
  - You're using one now.
- These machines are designed to facilitate the most common needs of the most people.
- However, the limits of these machines can be easily reached when performing modern neuroscience analysis.



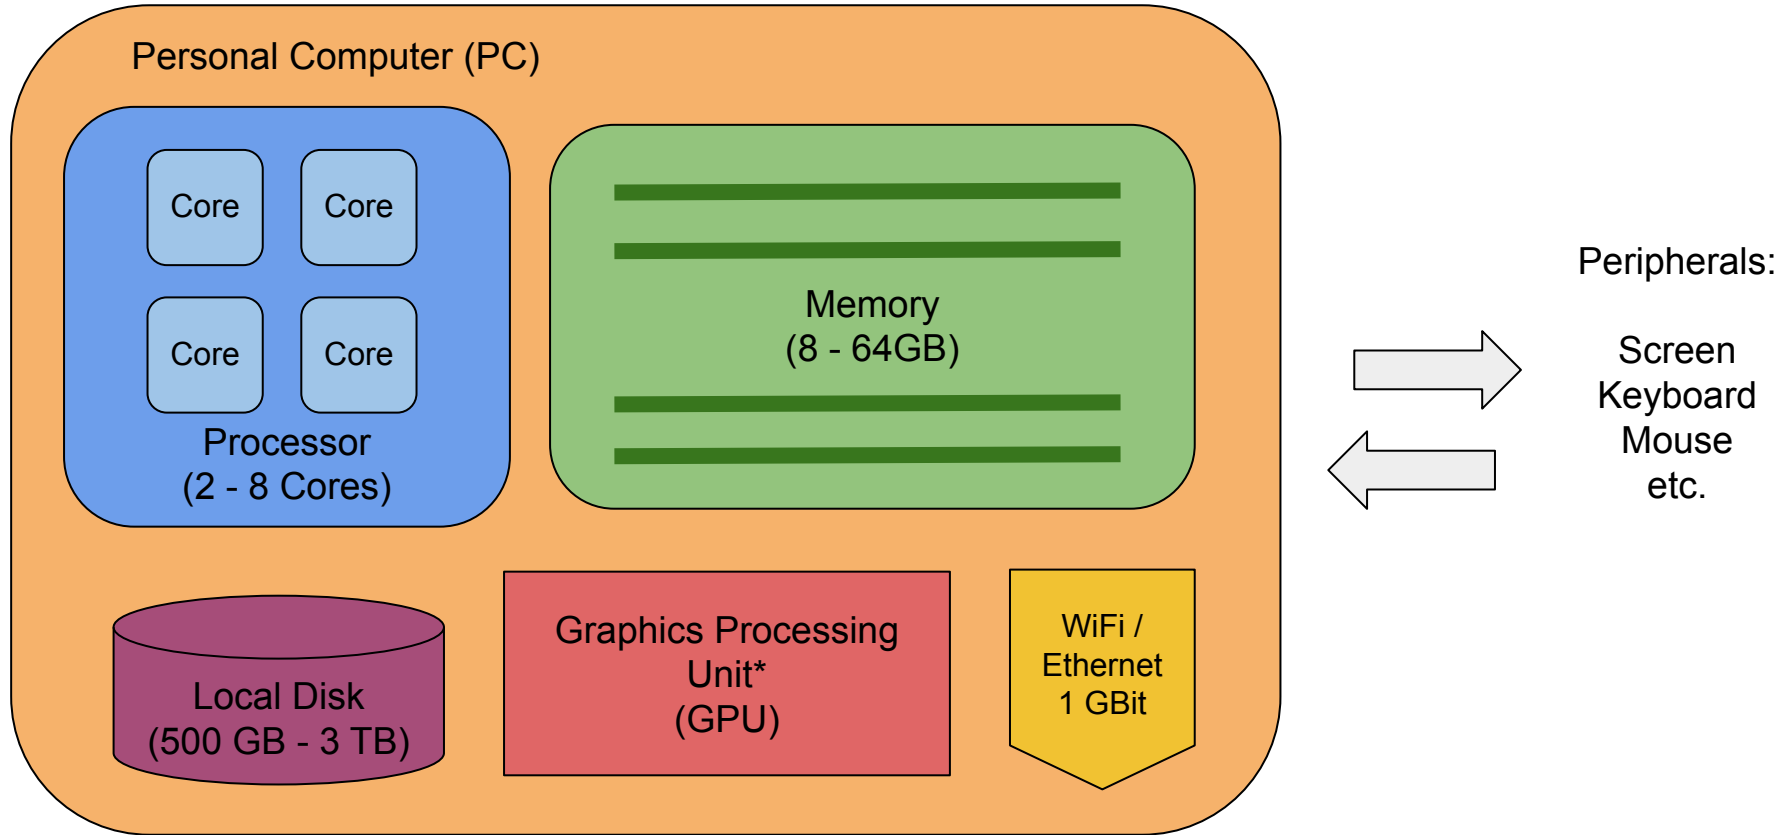
# Working on an HPC

- These are not “personal” systems, they are multi-user.
- You do not interact with them directly, but over a network connection.
- The interaction with your analysis entails:
  - Transferring data / results.
  - Managing your analysis “jobs”.
- There is minimal “interactive” computing performed on a HPC.
  - Typically little to no visualization is performed on HPC systems.
- They almost exclusively run Linux.

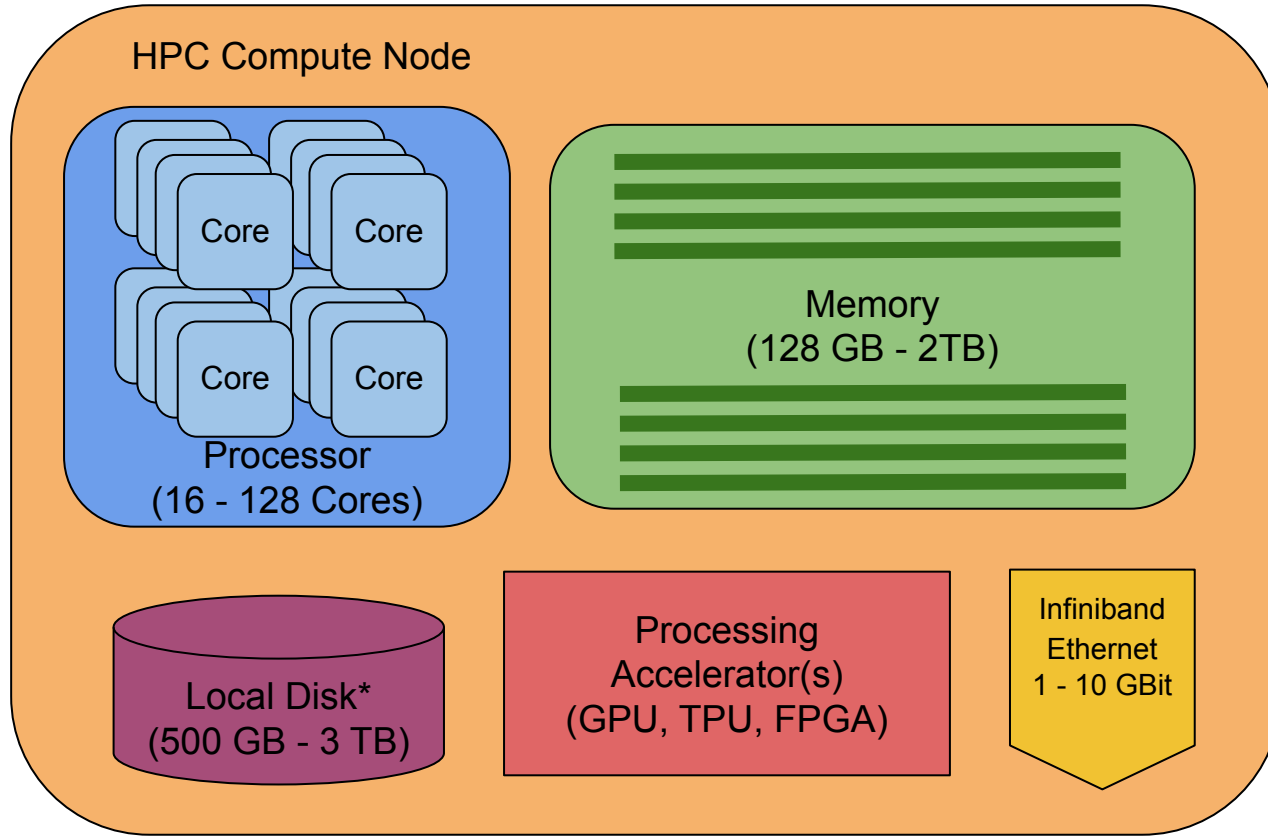


# The Schematic of a Computer





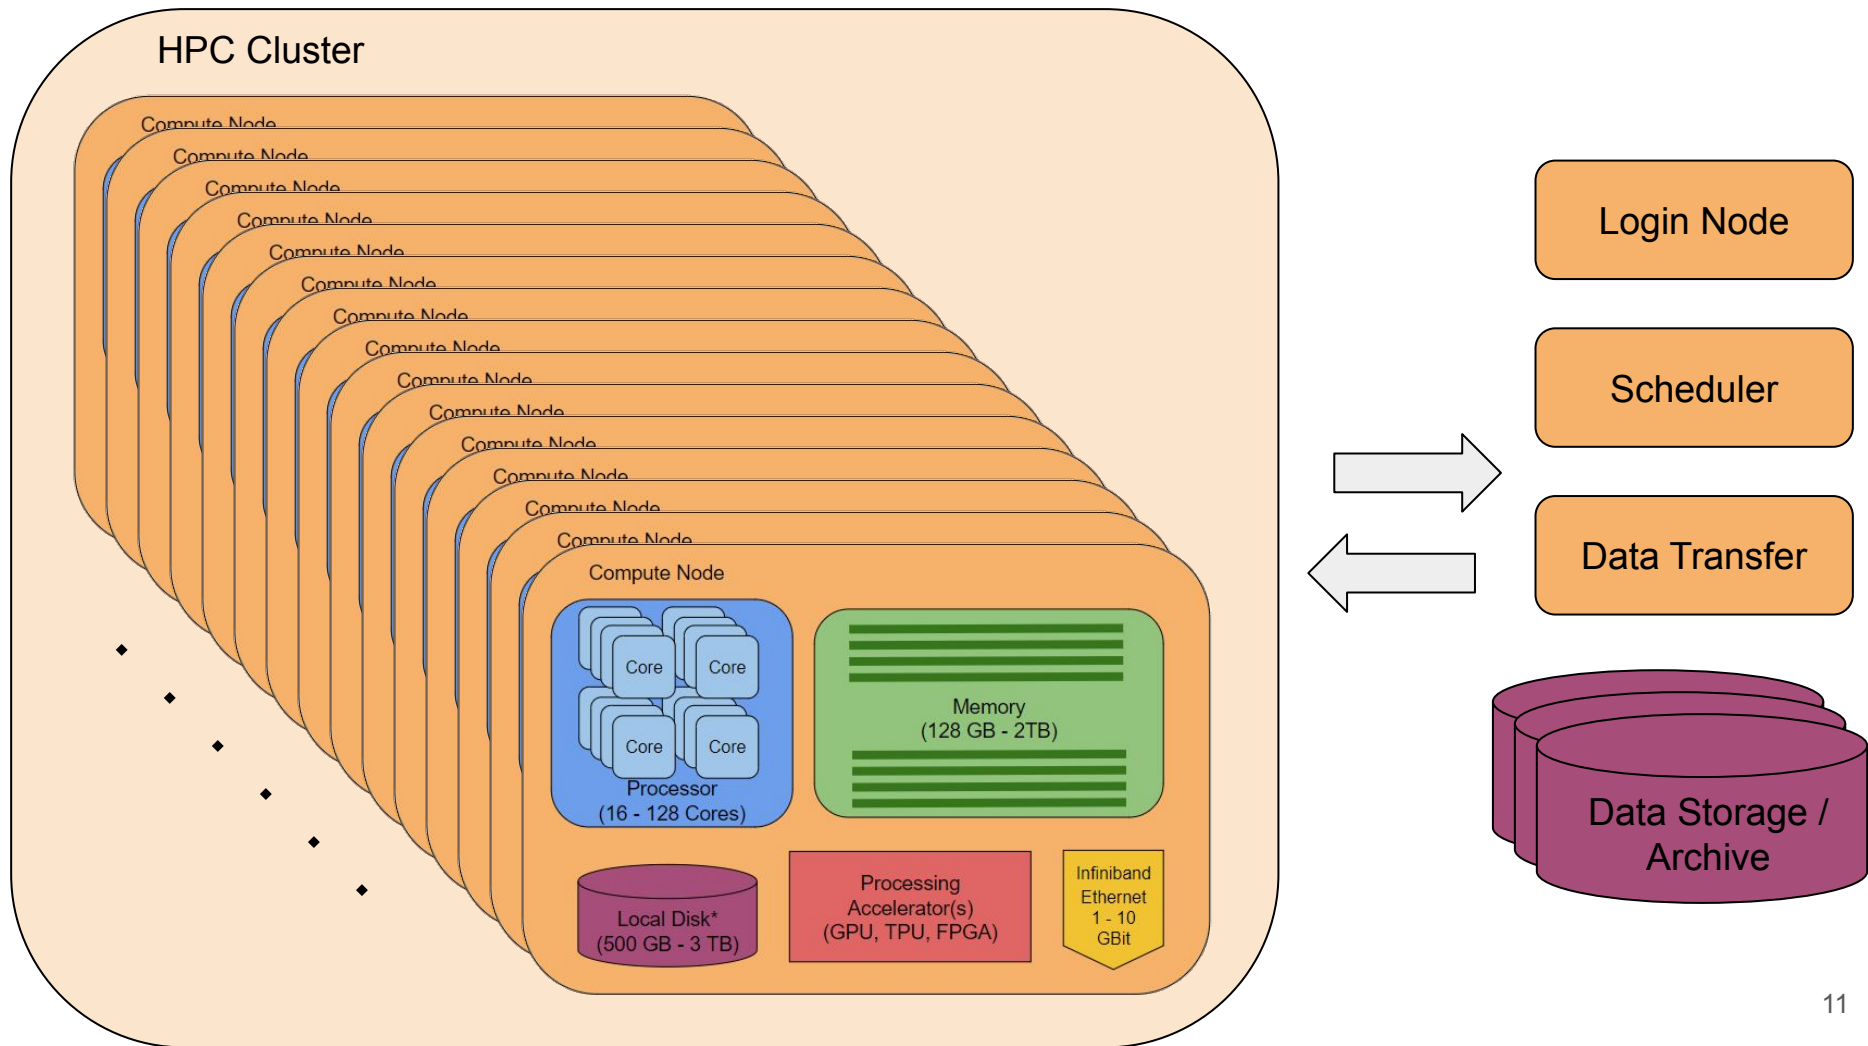
\* not necessarily present in all systems (laptops)



Peripherals:  
None

Connects to  
the rest of the  
compute  
cluster.

\* it may be less and / or have higher access speed



# Other important things to keep in mind

- For a single, small process, an HPC may be slower than your laptop
  - The clock speed of processors is often slower when more cores are present.
  - There is also the overhead of moving / staging data, waiting for walltime, etc.
- Some HPCs have a different system architecture
  - Most PCs are 64-bit (x86) system architectures.
    - ARM is also common for some laptops (Apple M1)
  - Some HPCs do not have an x86 architecture
    - Big Red II (Cray), Longhorn (PPCx64)
  - This may cause challenges for deploying your code.

# The Anatomy of a Computer

## PCs

- 1 CPU
  - 2 - 8 cores
  - 3.5 - 4.5 GHz
- 8 - 64 GB RAM
- 500 GB - 3 TB Storage
- 1 GPU
  - This may be built into the CPU
- WiFi / 1 GBit Ethernet
- x86 (maybe ARM) architecture
- Single-User System
- Good for most general usage
  - This is likely where you will develop your analysis.

## HPCs

- Multiple CPUs
  - 16 - 128 cores per CPU
  - 2.4 - 3.5 GHz
- 128 GB - 2 TB RAM
- 500 GB - 3 TB Storage
  - Not for long term storage
- 0 - 4 Processor Accelerators
  - GPUs w/ double precision CUDA
  - TPUs, FPGAs, Coprocessors, etc.
- 1 - 10 GBit Ethernet, Infiniband
- Different architectures (x86, PPC, Cray)
- Multi-User System
- Good for high throughput / large models.

# Working on the Cloud - An expensive middle ground

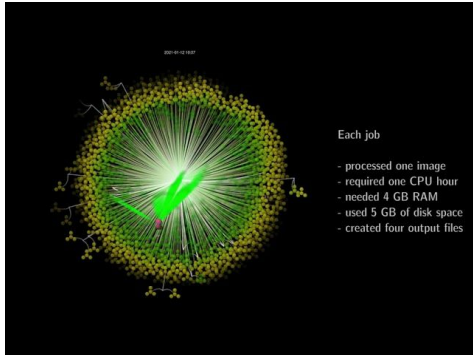
- This is most often an option through dedicated research proposals.
- The cloud system can be tailored to fit your specific analysis and data requirements.
  - A large virtual analysis computer
  - Your own dedicated cluster
- The data may be distributed on the cloud already.
  - Many open datasets are already on AWS.
- This will be *very* expensive for long term use.




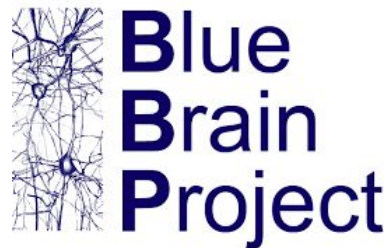
# The advantages of a HPC over a PC

- The ability to handle large datasets in a more time efficient manner.
- The option to “**Scale Out**” or “**Scale Up**” an analysis.
  - **Scale Out**: the ability to analyze independent parts of a dataset simultaneously
    - Independent permutations / cross-validations / simulations.
    - Different subjects through the same preprocessing preparations.
      - “Embarrassingly Parallel”
      - High Throughput Computing
  - **Scale Up**: create a larger single instance of computing resources to run a larger model, estimation, or analysis.
- We’ll be introducing you to “**Scaling Out**” your analysis in the example today.

# Scaling Out



# Scaling Up

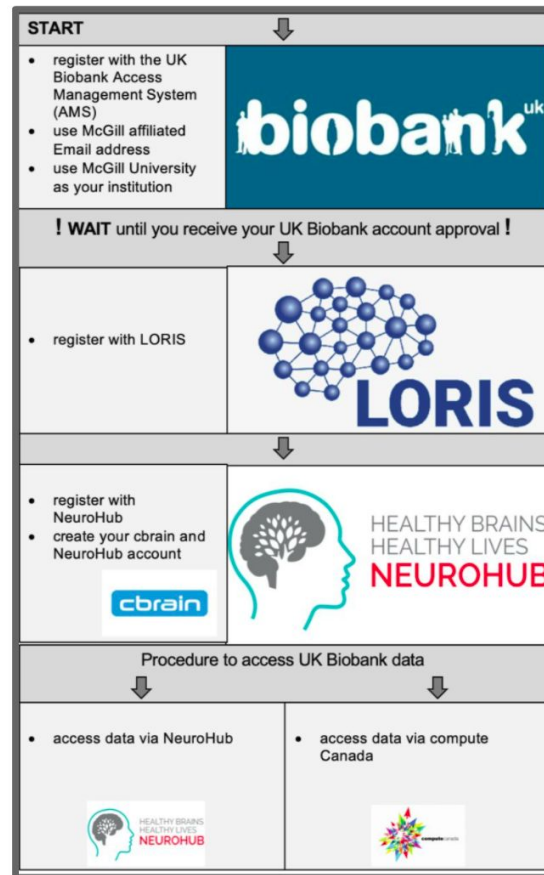




# UKBB Data Access

## Multi-step process!

1. Register with UKBB AMS
  - a. Project and data management platform
2. Register with LORIS
  - a. Project and data management platform
3. Get a NeuroHub account
  - a. Data and compute management platform
4. Get a Compute Canada account
  - a. High performance compute cluster



# Beginning to use a Cluster Computer

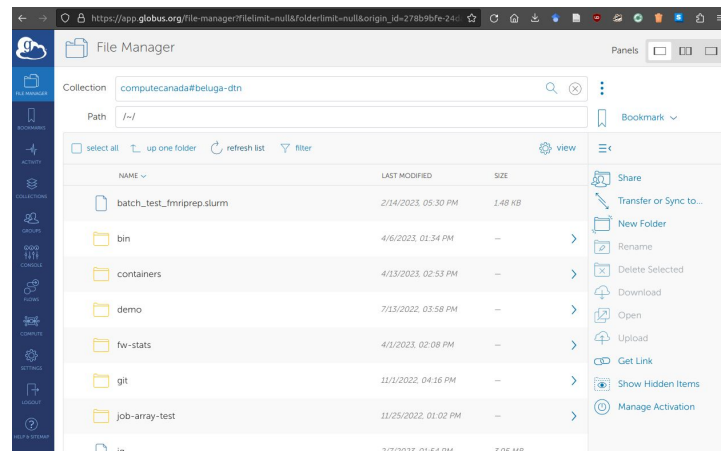
# Accessing the Cluster

- This is done exclusively through a remote connection.
  - `ssh`, a remote desktop client, etc.
- You will connect to the **Login Node** of the cluster.
  - The login in node manages the coordination of user processes on the cluster
  - It also reports the overall status of the system.
- You will submit your jobs and query their status from the Login Node.
  - DO NOT run analyses or do other work on the Login Node.



# Getting your data on the cluster

- I assume your data is on a computer not accessible by the cluster.
- You will use a file transfer tool to move your data from where it is to the data server.
  - scp, sftp, rsync, FileZilla, mobaXterm, etc.
  - This data server will be accessible to the compute nodes
- This may be done through the Login Node, but there may be a specific Data Transfer Node to move data.
  - Check with the system administrators.



# Getting your data on the cluster - Things to Remember

- Be aware of the policies of this server!
  - Files are often automatically removed after a (potentially short) window of inactivity.
  - Your analysis code, jobs, and data are probably not backed up on the HPC.
- It is your responsibility to make sure your data and code are securely backed up somewhere else.
- Once your analysis is complete you need to have a plan for storing your data long term.
  - OSF
  - OpenNeuro

# Getting your software on the cluster

- The cluster administrators will likely need to install it for you.
  - If you are confident you know how the tool works, you may be able to set it up yourself.
- Many tools will be readily available on the cluster through **modules**.
  - These are an efficient way to add and remove tools installed on the cluster to your session.
  - They are like an interactive `venv` for any tool.
  - They also allow for easier management of different versions of the same tool.
  - Hopefully, most of the tools you need will already be installed.
- Running your code in **containers** guarantees you have the tools you need to run your analysis.
  - You do have to have (or be able to build) a working container of the tool.

# Getting your software on the cluster - Modules

- How most software is made available to users on a HPC.
- A convenient way to set up basic development environments and change between versions.

```
> module avail  
> module load python  
> module load python/3.10.2  
> module unload python  
> module swap python/3.9.6
```





```
[bcmcphe@beluga4 ~]$ module avail
```

```
----- Cluster specific modules -----  
httpproxy/1.0
```

```
----- MPI-dependent avx512 modules -----  
abinit/9.2.2      (chem)    etsf_io-mpi/1.0.4      mscg/1.7.3.1      (chem)    parsplce/1.1      repasthpc/2.3.0      (bio,D)  
abinit/9.6.2      (chem,D)   fds/6.7.5             mumps-metis/5.2.1  (t)         petsc/3.12.4      (t)         rosetta/3.10        (chem)  
adol-c/2.7.2      (chem)    fds/6.7.6             mumps-parmetis/5.3.5 (t)         petsc/3.13.6      (t)         rosetta/3.12        (chem)  
apbs/1.3          (chem)    fds/6.7.7             ncl/6.6.2         (vis)        petsc/3.14.1      (t)         rosetta/2019.21.60746 (chem,D)  
berkeleygw/2.1.0  (phys)    fds/6.7.8             ncview/2.1.8      (vis)        petsc/3.15.0      (t,D)       scalapack/2.1.0     (math)  
berkeleygw/3.0.1  (phys,D)  fftw-mpi/3.3.8        nektar++/5.0.1     (math)       pfft/1.0.8-alpha  (math)       scotch/6.0.9       (math)  
bigdft/1.8.3      (chem)    ga/5.7.2              netcdf-c++-mpi/4.2 (io)         phylobayes-mpi/20180420 (bio)       shengbte/1.1.1     (phys)  
boost-mpi/1.72.0  (t)       globalarrays/5.7.2    netcdf-c++4-mpi/4.3.1 (io)         phylobayes-mpi/20201026 (bio,D)     siesta/4.0.1       (chem)  
cdo/1.9.8         (geo)     globalarrays/5.8      netcdf-fortran-mpi/4.5.2 (io)         plumed/2.6.1      (chem)       siesta/4.1-b4       (chem)  
cdo/2.0.4         (geo)     glost/0.3.1           netcdf-mpi/4.7.4   (io)         plumed/2.6.2      (chem)       siesta/4.1-MaX-3.0  (chem)  
cdo/2.0.5         (geo,D)   hdf5-mpi/1.10.6       neuron/7.8.2       (bio)        plumed/2.7.0      (chem)       siesta/4.1.5        (chem,D)  
cfour-mpi/2.1     (chem)    hdf5-mpi/1.12.1       neuron/8.0.0       (bio,D)      plumed/2.7.1      (chem)       slepc/3.14.2        (chem)  
cgns/3.4.1        (phys)    hpl/2.3               nwchem/6.8.1      (chem)       plumed/2.7.3      (chem,D)     sundials/2.7.0      (chem)  
cgns/4.1.0        (phys)    hypre/2.20.0          nwchem/7.0.2-p1    (chem,D)     pnetcdf/1.9.0     (io)         sundials/5.3.0     (D)  
cgns/4.1.2        (phys,D)  ima3/20210120         octopus/10.1       (chem)       pnetcdf/1.10.0    (io)         wannier90-abinit/2.0.1.1 (chem)  
comblas/1.6.2     (chem)    lammps-omp/20201029   openfoam-extend/4.1 (phys)       pnetcdf/1.12.2    (io,D)       wannier90/3.1.0     (chem)  
cp2k/7.1          (chem)    lammps-omp/20210929   openfoam/6         (phys)       psi4/1.3.2        (chem)       wps/4.1            (geo)  
cp2k/8.2          (chem)    latte/1.2.1           openmolcas/20.10   (chem)       psi4/1.4          (chem,D)     wps/4.2            (geo,D)  
cp2k/9.1          (chem,D)  libcf/1.0.3           openmx/3.9         (chem)       quantumpresso/6.5  (chem)       wrf/4.1.3          (geo)  
cpmd/4.3          (chem)    libgridxc-mpi/0.8.0   openmx/3.9.9       (D)          quantumpresso/6.6  (chem)       wrf/4.2.1          (geo)  
cslib/20180813    (t)       met/9.1.1             osu-micro-benchmarks/5.6.2 (t)         quantumpresso/6.7  (chem)       wrf/4.3.3          (geo,D)  
dakota/6.13       (t)       mpas/7.0              p4est/2.2         (math)       quantumpresso/6.8  (chem)       yambo/5.0.4        (chem)  
delft3d/62441     (chem)    mpi4py/3.0.3          parallelio/2.5.4   (t)          quantumpresso/7.0  (chem,D)     yast/0.9.0         (t)  
dl_poly4/4.10.0   (chem)    mpi4py/3.1.2          paraview-offscreen/5.8.0 (vis)        raxml/8.2.12      (bio)  
elpa/2020.05.001  (math)    mpi4py/3.1.3          parmetis/4.0.3     (math)       ray/3.0.1         (bio)  
esmf/8.0.1        (geo)     mrbayes/3.2.7         parmgridgen/1.0    (math)       repasthpc/2.2.0    (bio)
```

# Running your analysis - Working with a Scheduler

- With your data and `env` ready you can run your analysis.
- You will do this by submitting your work as a **Job**.
- **Jobs** are submitted to the Scheduler to be run.
  - This deploys the analysis onto the compute nodes.
- Your entire analysis will need to be scripted.
  - You cannot interact with or manually input information to a job when it's running.
- Once your analysis script is complete, you create your job script for the analysis.
  - This script describes the resources needed to successfully run your analysis on the cluster.



# Running your analysis - Available Resources / Queues

- You will need to submit your work to a compute cluster that can accommodate your resource request.
- Specifically, this means you may have to submit your jobs to the GPU queue if your analysis requires a GPU to execute.
  - This will be in addition to requesting the GPU cores / memory
- Similar queues exist for jobs that are:
  - Expected to run for a long time ( > 24 hours)
  - High memory requirement ( > 256GB)
- The queues available on a cluster reflect the hardware available.
  - Often you will select the cluster you will work on based on the queues you plan to use.

# Running your analysis - Creating Job Scripts

- Your job scripts define the system resources required by your analysis.
- The job script will be submitted to the scheduler from the login node to be added to the queue.
  - The scheduler uses your resource requests and your priority to determine what runs when.
- Your outputs, along with logs, will be created for each job you submit.
- The example today will assume you are **Scaling Out** your analysis.
  - You want to run a common process across many subjects.
  - i.e. fMRIPrep, FreeSurfer, etc.

demo\_subj01.slurm

```
#!/bin/bash

#SBATCH --job-name=demo_subj01      # job name
#SBATCH --nodes=1                   # run on a single node
#SBATCH --ntasks=1                  # run on a single CPU
#SBATCH --cpus-per-task=1           # run on a single core
#SBATCH --mem=1gb                   # job memory request
#SBATCH --time=00:05:00             # time limit hrs:min:sec
#SBATCH --error=./jobs/logs/demo_subj01_%j.err # standard error from job
#SBATCH --output=./jobs/logs/demo_subj01_%j.out # standard output from job
#SBATCH --account=def-jbpoline      # define your affiliation

## ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
## MODIFY THE RESOURCE REQUESTS ABOVE FOR YOUR JOB

## MODIFY THE CODE BELOW TO CALL YOUR ANALYSIS
## vvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvv

## your job script generalized to a subject ID
bash ./analysis.sh subj01
```

# Commands for working with Jobs

- **sbatch**
  - Submit a job to the scheduler.
  - `> sbatch job.slurm`
- **squeue**
  - Monitor the status of jobs on the cluster.
  - `> squeue -u USERNAME`
- **scancel**
  - Cancels a job that hasn't completed.
  - `> scancel JOBID`
- **sacct**
  - Provides a summary of the jobs your have recently submitted.
  - `> sacct`
- **srun**
  - Start an interactive job to test that your job will work correctly.
  - `> srun --nodes=1 --ntasks-per-node=1 --time=01:00:00 --pty bash -i`



There are a lot of resources available for learning about HPCs online and in your lab. Find a process that works for you and ASK how others have solved similar problems!

# Running your analyses - be efficient

- For every job you run, you will have:
  - The input data - *the same files for each subject*
  - The analysis script - *the same analysis for each subject*
  - The job script and logs - *a separate job for every subject with logs for each attempt*
  - The results - *whatever output your analysis produces that you intend to keep*
- Just like your analyses, you should script as much of your cluster work as possible.
  - You should never be individually making or modifying job files for subjects.
  - Having scripts for moving and backing up data is a good idea, too.



# The Example

- The example today reports back a short log about the system status of the node before waiting.
  - This will let you see the different nodes your jobs ran on.
- The analysis input wants a subject ID as an input
  - This is is likely how you would set up a symmetric analysis across your subjects.
- The goal is to provide a basic working template for your job scripts that you can modify.

```
bcmcphe@ThinkPad-T14s:~/git/qis612-2023/Lectures/12-High_Performance_Computing$ ls -l  
total 24  
-rwxr-x--- 1 bcmcpher bcmcpher 645 May 5 02:47 00_copy_data.sh  
-rwxr-x--- 1 bcmcpher bcmcpher 107 May 5 02:47 01_create_jobs.sh  
-rwxr-x--- 1 bcmcpher bcmcpher 81 May 5 02:47 02_submit_jobs.sh  
-rwxr-x--- 1 bcmcpher bcmcpher 42 May 5 02:47 03_inspect_results.sh  
-rwxr-x--- 1 bcmcpher bcmcpher 653 May 5 02:47 zz_analysis.sh  
-rwxr-x--- 1 bcmcpher bcmcpher 1214 May 5 02:47 zz_mk_job.sh  
bcmcphe@ThinkPad-T14s:~/git/qis612-2023/Lectures/12-High_Performance_Computing/hpc-example/scripts$ cat 01_create_jobs.sh  
#!/bin/bash  
  
find ../data -type d -name 'subj*' -printf '%f\n' | sort | xargs -n 1 -I {} ./zz_mk_job.sh {}  
bcmcphe@ThinkPad-T14s:~/git/qis612-2023/Lectures/12-High_Performance_Computing/hpc-example/scripts$ cat 02_submit_jobs.sh  
#!/bin/bash  
  
find jobs -type f -name '*.slurm' | sort | xargs -n 1 -I {} sbatch {}  
bcmcphe@ThinkPad-T14s:~/git/qis612-2023/Lectures/12-High_Performance_Computing/hpc-example/scripts$ cat zz_mk_job.sh  
#!/bin/bash  
  
SUBJ=$1  
  
cat << EOF > ../jobs/demo_${SUBJ}.slurm  
#!/bin/bash  
  
#SBATCH --job-name=demo_${SUBJ} # job name  
#SBATCH --nodes=1 # run on a single node  
#SBATCH --ntasks=1 # run on a single CPU  
#SBATCH --cpus-per-task=1 # run on a single core  
#SBATCH --mem=1gb # job memory request  
#SBATCH --time=00:05:00 # time limit hrs:min:sec  
#SBATCH --error=./jobs/logs/demo_${SUBJ}_%.err # standard error from job  
#SBATCH --output=./jobs/logs/demo_${SUBJ}_%.out # standard output from job  
#SBATCH --account=def-jbpoline # define your affiliation  
  
## #####  
## MODIFY THE RESOURCE REQUESTS ABOVE FOR YOUR JOB  
  
## MODIFY THE CODE BELOW TO CALL YOUR ANALYSIS  
## vvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvv  
  
## your job script generalized to a subject ID  
bash ./zz_analysis.sh ${SUBJ}  
  
##  
## use a container  
##  
  
#module load aptainer  
  
## your job script generalized to a subject ID  
#aptainer exec -H /scratch/$USER/hpc-demo -B /scratch/$USER/hpc-demo ../container/ubuntu.sif bash ./zz_analysis.sh ${SUBJ}  
  
EOF
```

1. Connect to the HPC.
2. Move the example scripts to the HPC.
3. Make and submit your jobs.
4. Monitor the job status.
5. Cancel a job.
6. Run an interactive job.
7. Modify your resource requests / Explore documentation.
8. Run the script in a container.

# Demo Redux

## Step 0: Create a container

```
module load singularity
```

```
singularity build ../container/ubuntu.sif docker://ubuntu:latest
```

- From the login node, load the singularity module
  - Interactive sessions may not have internet access to download an image
- Call the build function pointing to the docker image you want to import
  - `docker://` - look on dockerhub
  - `Ubuntu:latest` for the most recently tagged ubuntu image
- This will create the singularity container you can use on Compute Canada



## Step 1.1 - Make a template job script w/ Singularity

- You will also be able to call your jobs with a singularity container.
- You can call multiple scripts or functions within your job.
- Don't forget to load the module!

```
#!/bin/bash

#SBATCH --job-name=demo_${SUBJ}           # job name
#SBATCH --nodes=1                         # run on a single node
#SBATCH --ntasks=1                       # run on a single CPU
#SBATCH --cpus-per-task=1                 # run on a single core
#SBATCH --mem=1gb                         # job memory request
#SBATCH --time=00:05:00                   # time limit hrs:min:sec
#SBATCH --error=./jobs/logs/demo_${SUBJ}_${j}.err # standard error from job
#SBATCH --output=./jobs/logs/demo_${SUBJ}_${j}.out # standard output from job
#SBATCH --account=def-jbpoline            # define your affiliation


## ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
## MODIFY THE RESOURCE REQUESTS ABOVE FOR YOUR JOB


## MODIFY THE CODE BELOW TO CALL YOUR ANALYSIS
## vvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvvv

module load apptainer

apptainer exec -H /scratch/$USER/hpc-demo -B /scratch/$USER/hpc-demo ../container/ubuntu.sif bash ./zz_analysis.sh ${SUBJ}
```

## Step 1.2 - Make a job file for each subject

```
find ../data -type d -name 'subj*' -printf '%f\n' | sort | xargs -n 1 -I {} ./zz_mk_job.sh {}
```

- `find ../data`
  - Find in the data folder
- `-type d -name *.slurm`
  - All directories named `subj*`
- `-printf '%f\n'`
  - Only print the directory name. This is each subject ID for this study.
- `| sort | xargs -n 1 -I {} zz_mk_job.sh {}`
  - Then sort the subject IDs and pass each ID as an input to `zz_mk_job.sh` to create a job script from your template for each subject.



## Step 2 - Submit your job scripts

```
find jobs -type f -name *.slurm | sort | xargs -n 1 -I {} sbatch {}
```

- `find jobs`
  - Find in the jobs folder
- `-type f -name *.slurm`
  - All files named \*.slurm
- `| sort`
  - Sort them in ascending order
- `| xargs -n 1 I {} sbatch {}`
  - Then pass each file as an input to `sbatch` to launch the job

## Step 3 - Extract your results

- Once the jobs are done, you can extract your results
  - This may mean moving the data off the compute cluster to visualize
- For the tutorial your outputs are text. So you can

- `grep -r 'error' ./logs/*`

- To look for errors (or other messages) in the logs

- `grep -r 'compute' ../results/*`

- To look for the result the “analysis” returned.

## Step 4 - Make it your own!

- The steps outlined here are a good starting point of some best practices, but there is so much you can do with HPCs.
  - There are many different ways to effectively use these tools and manage your jobs.
    - Job Arrays, for example
  - This is a wide domain to explore
- Learn what you can from others and try your best to understand these kinds of systems, they'll be an invaluable tool for the rest of your career.

Thanks